

Transfer Learning with

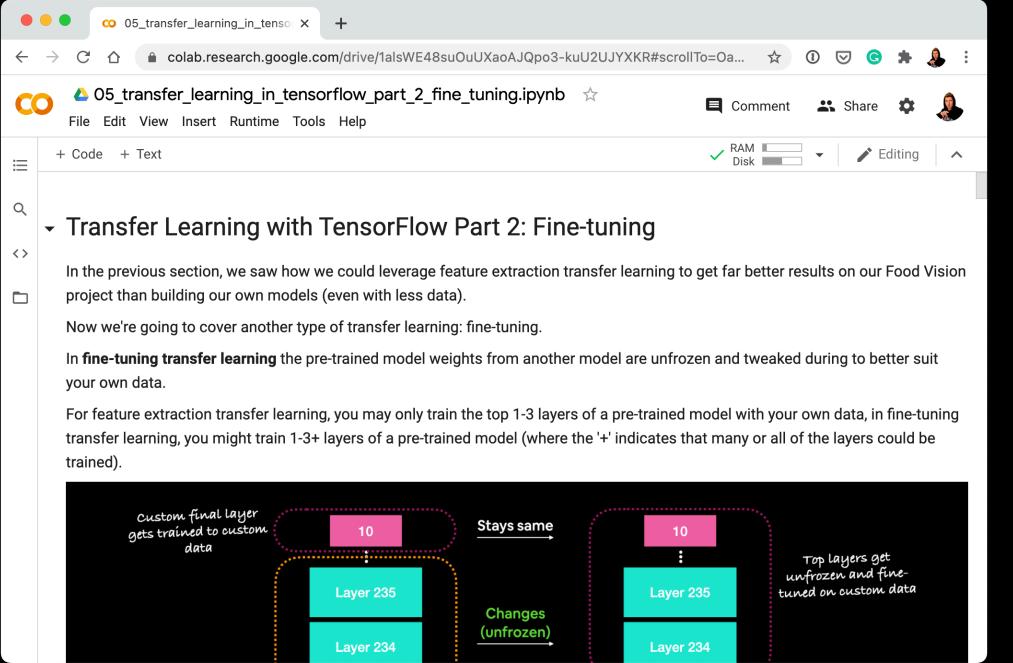


TensorFlow

Part 2: Fine-tuning

Where can you get help?

- Follow along with the code



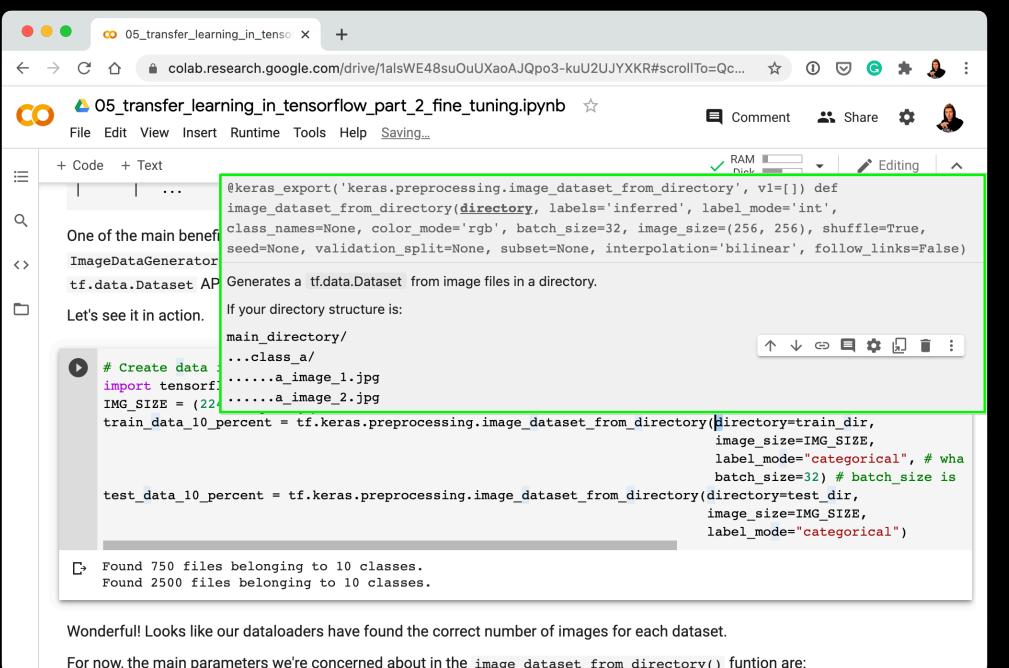
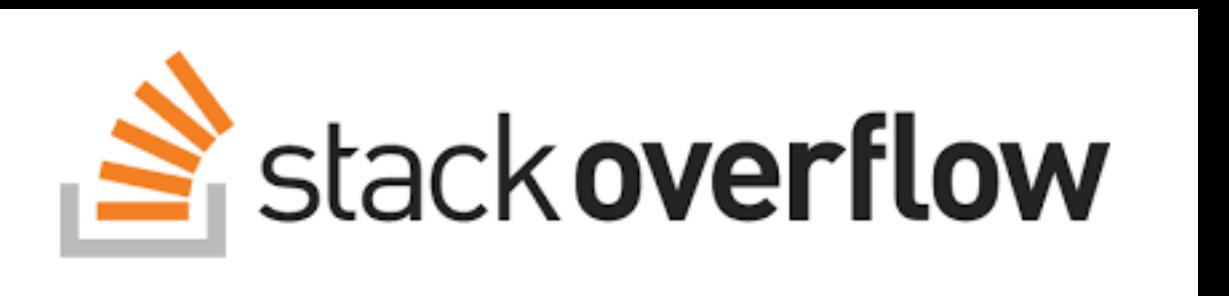
"If in doubt, run the code"

- Try it for yourself

- Press SHIFT + CMD + SPACE to read the docstring



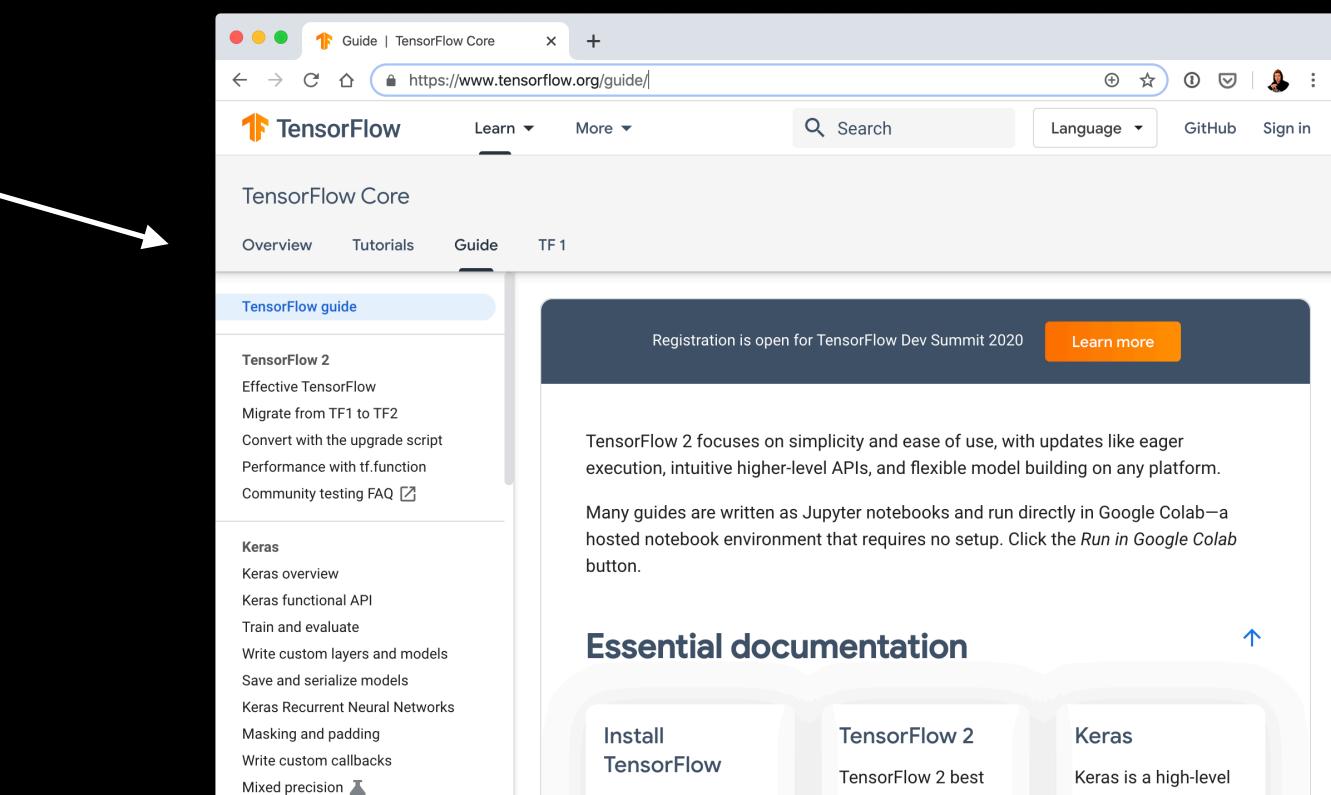
- Search for it



- Try again

- Ask (don't forget the Discord chat!)

(yes, including the "dumb" questions)

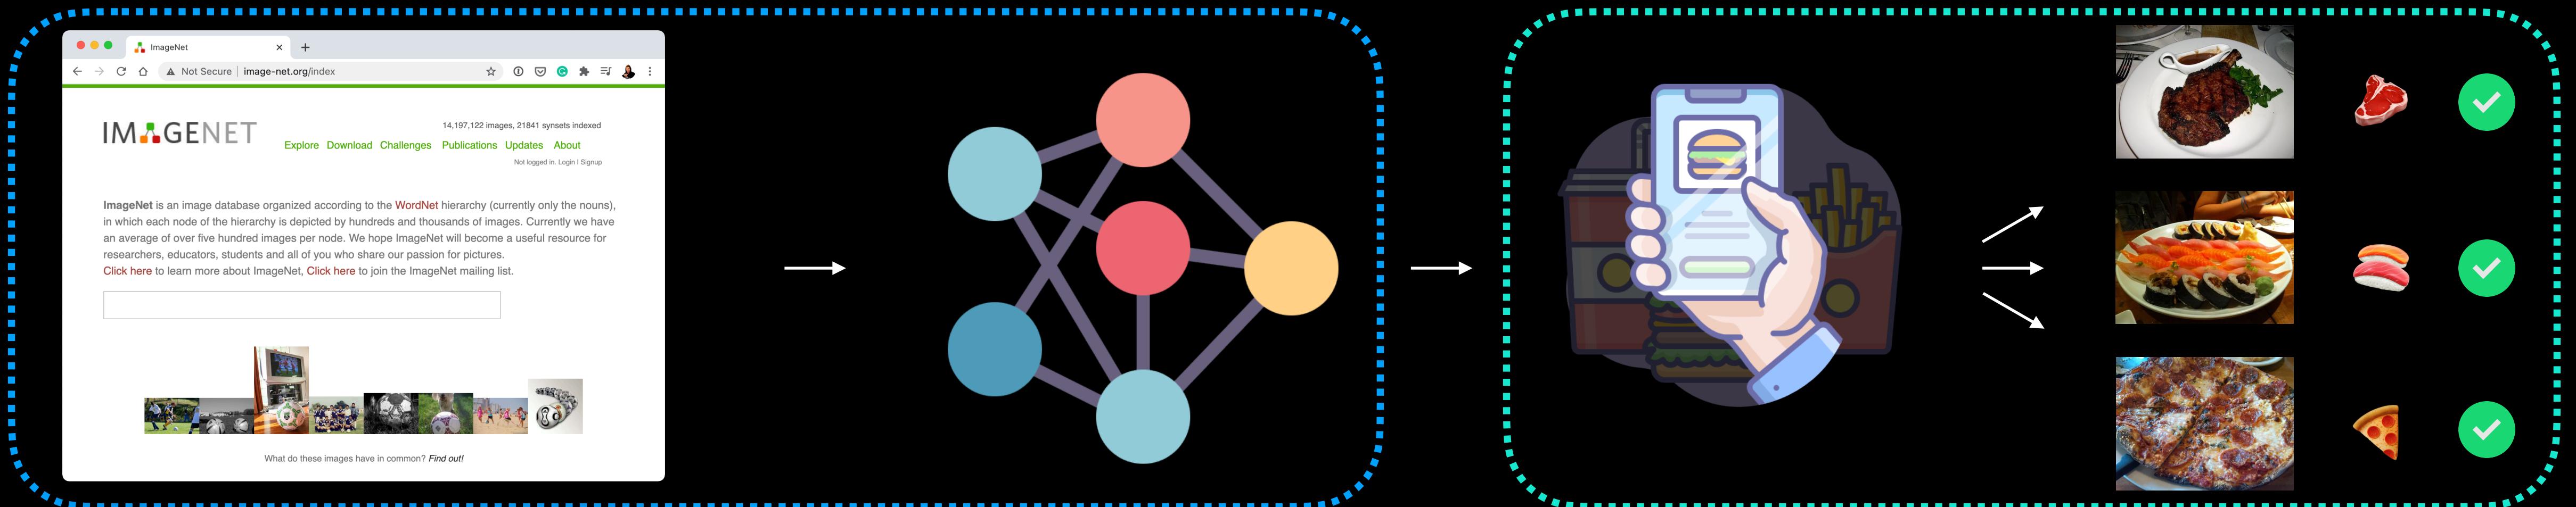


“What is transfer learning?”

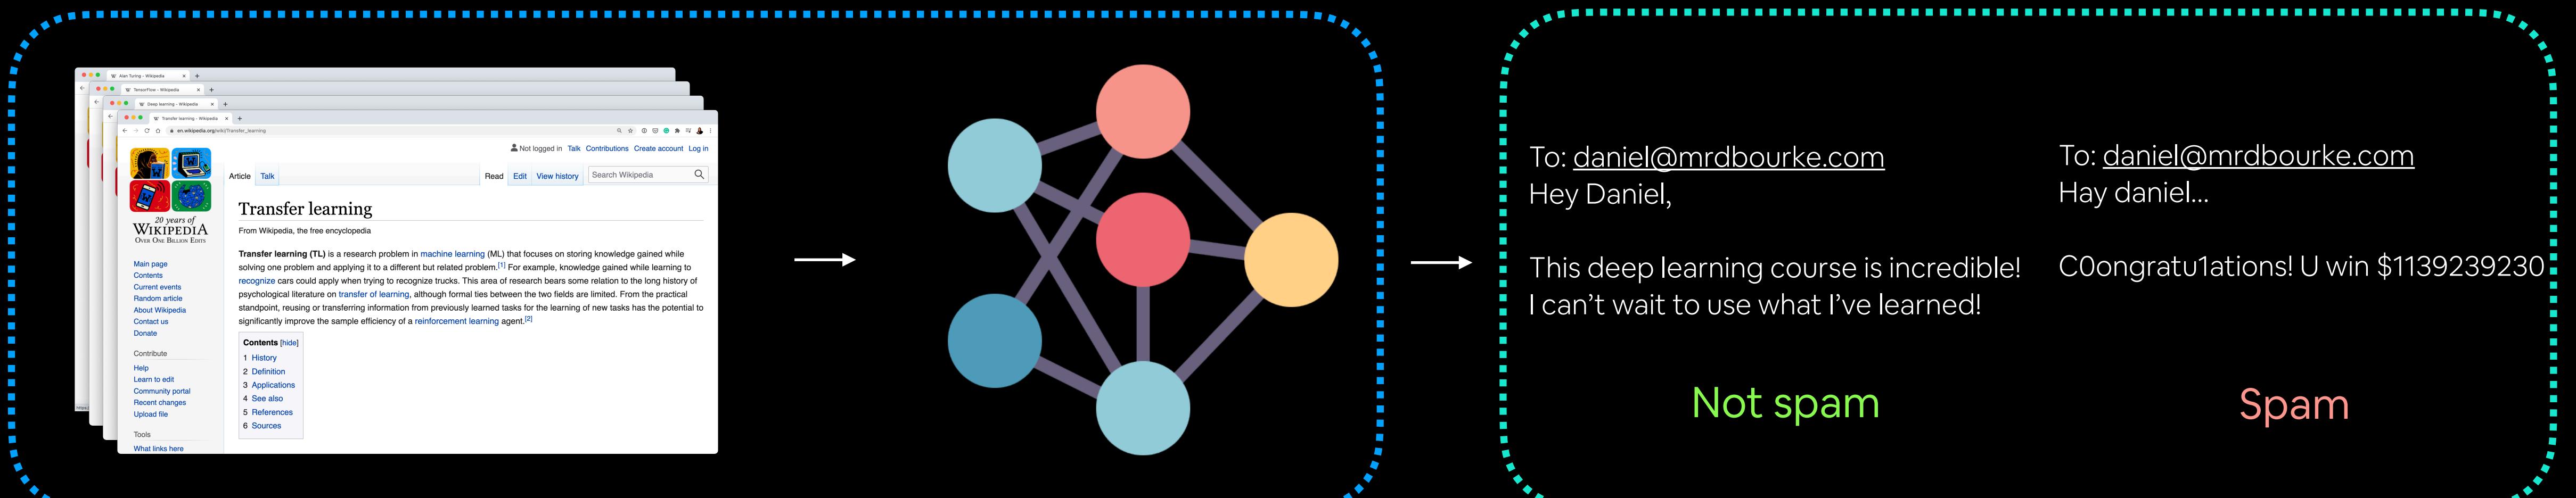
Surely someone has spent the time crafting the right model for the job...

Example transfer learning use cases

Computer vision



Natural language processing



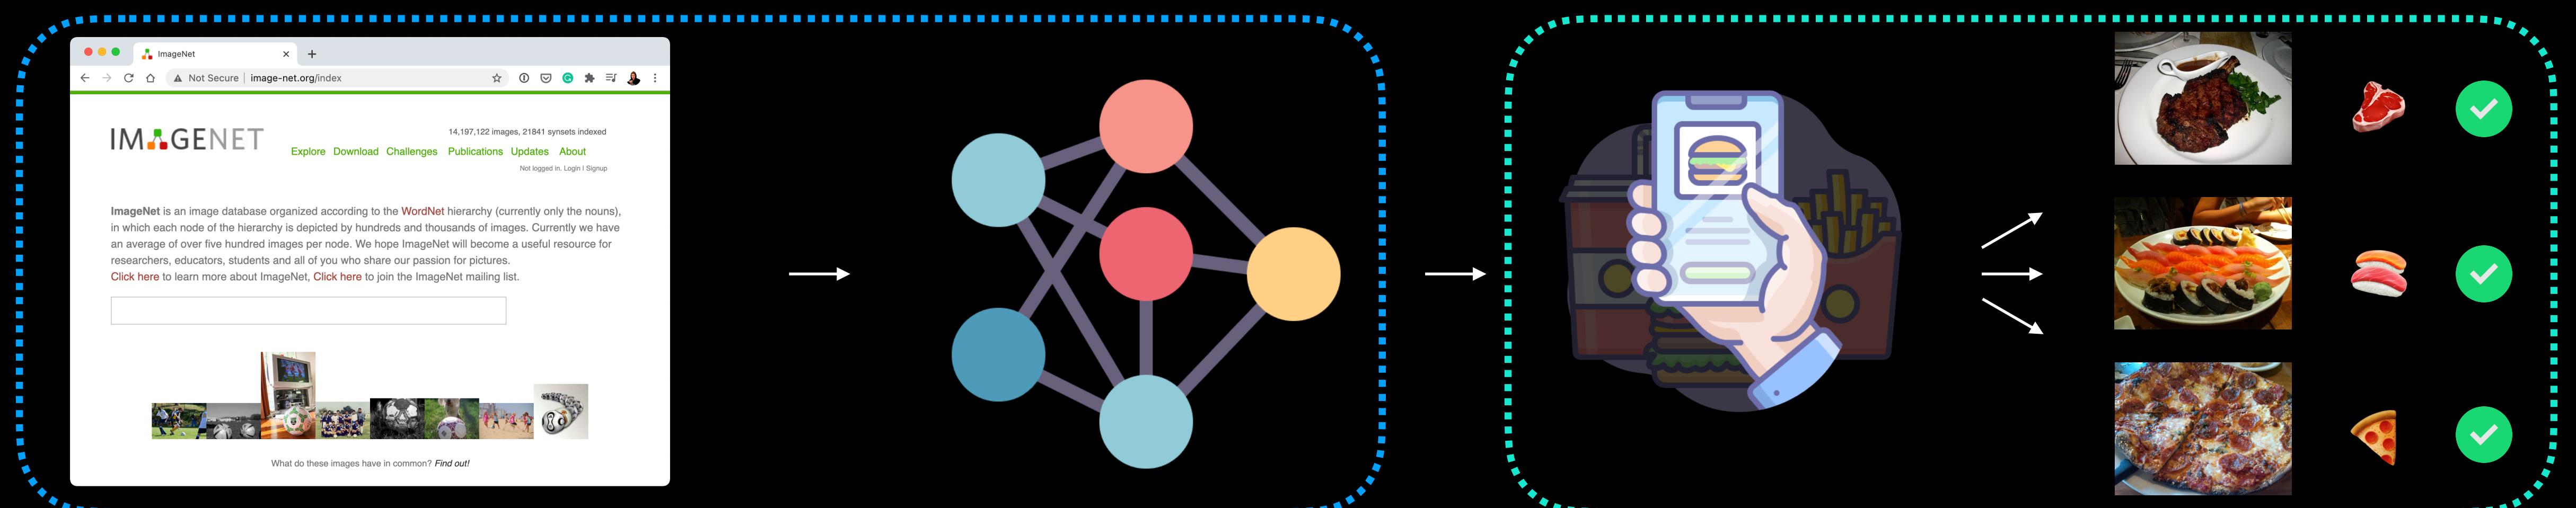
Model learns patterns/weights from similar problem space

Patterns get used/tuned to specific problem

“Why use transfer learning?”

Why use transfer learning?

- Can leverage an existing neural network architecture **proven to work** on problems similar to our own
- Can leverage a working network architecture which has **already learned patterns** on similar data to our own (often results in great results with less data)



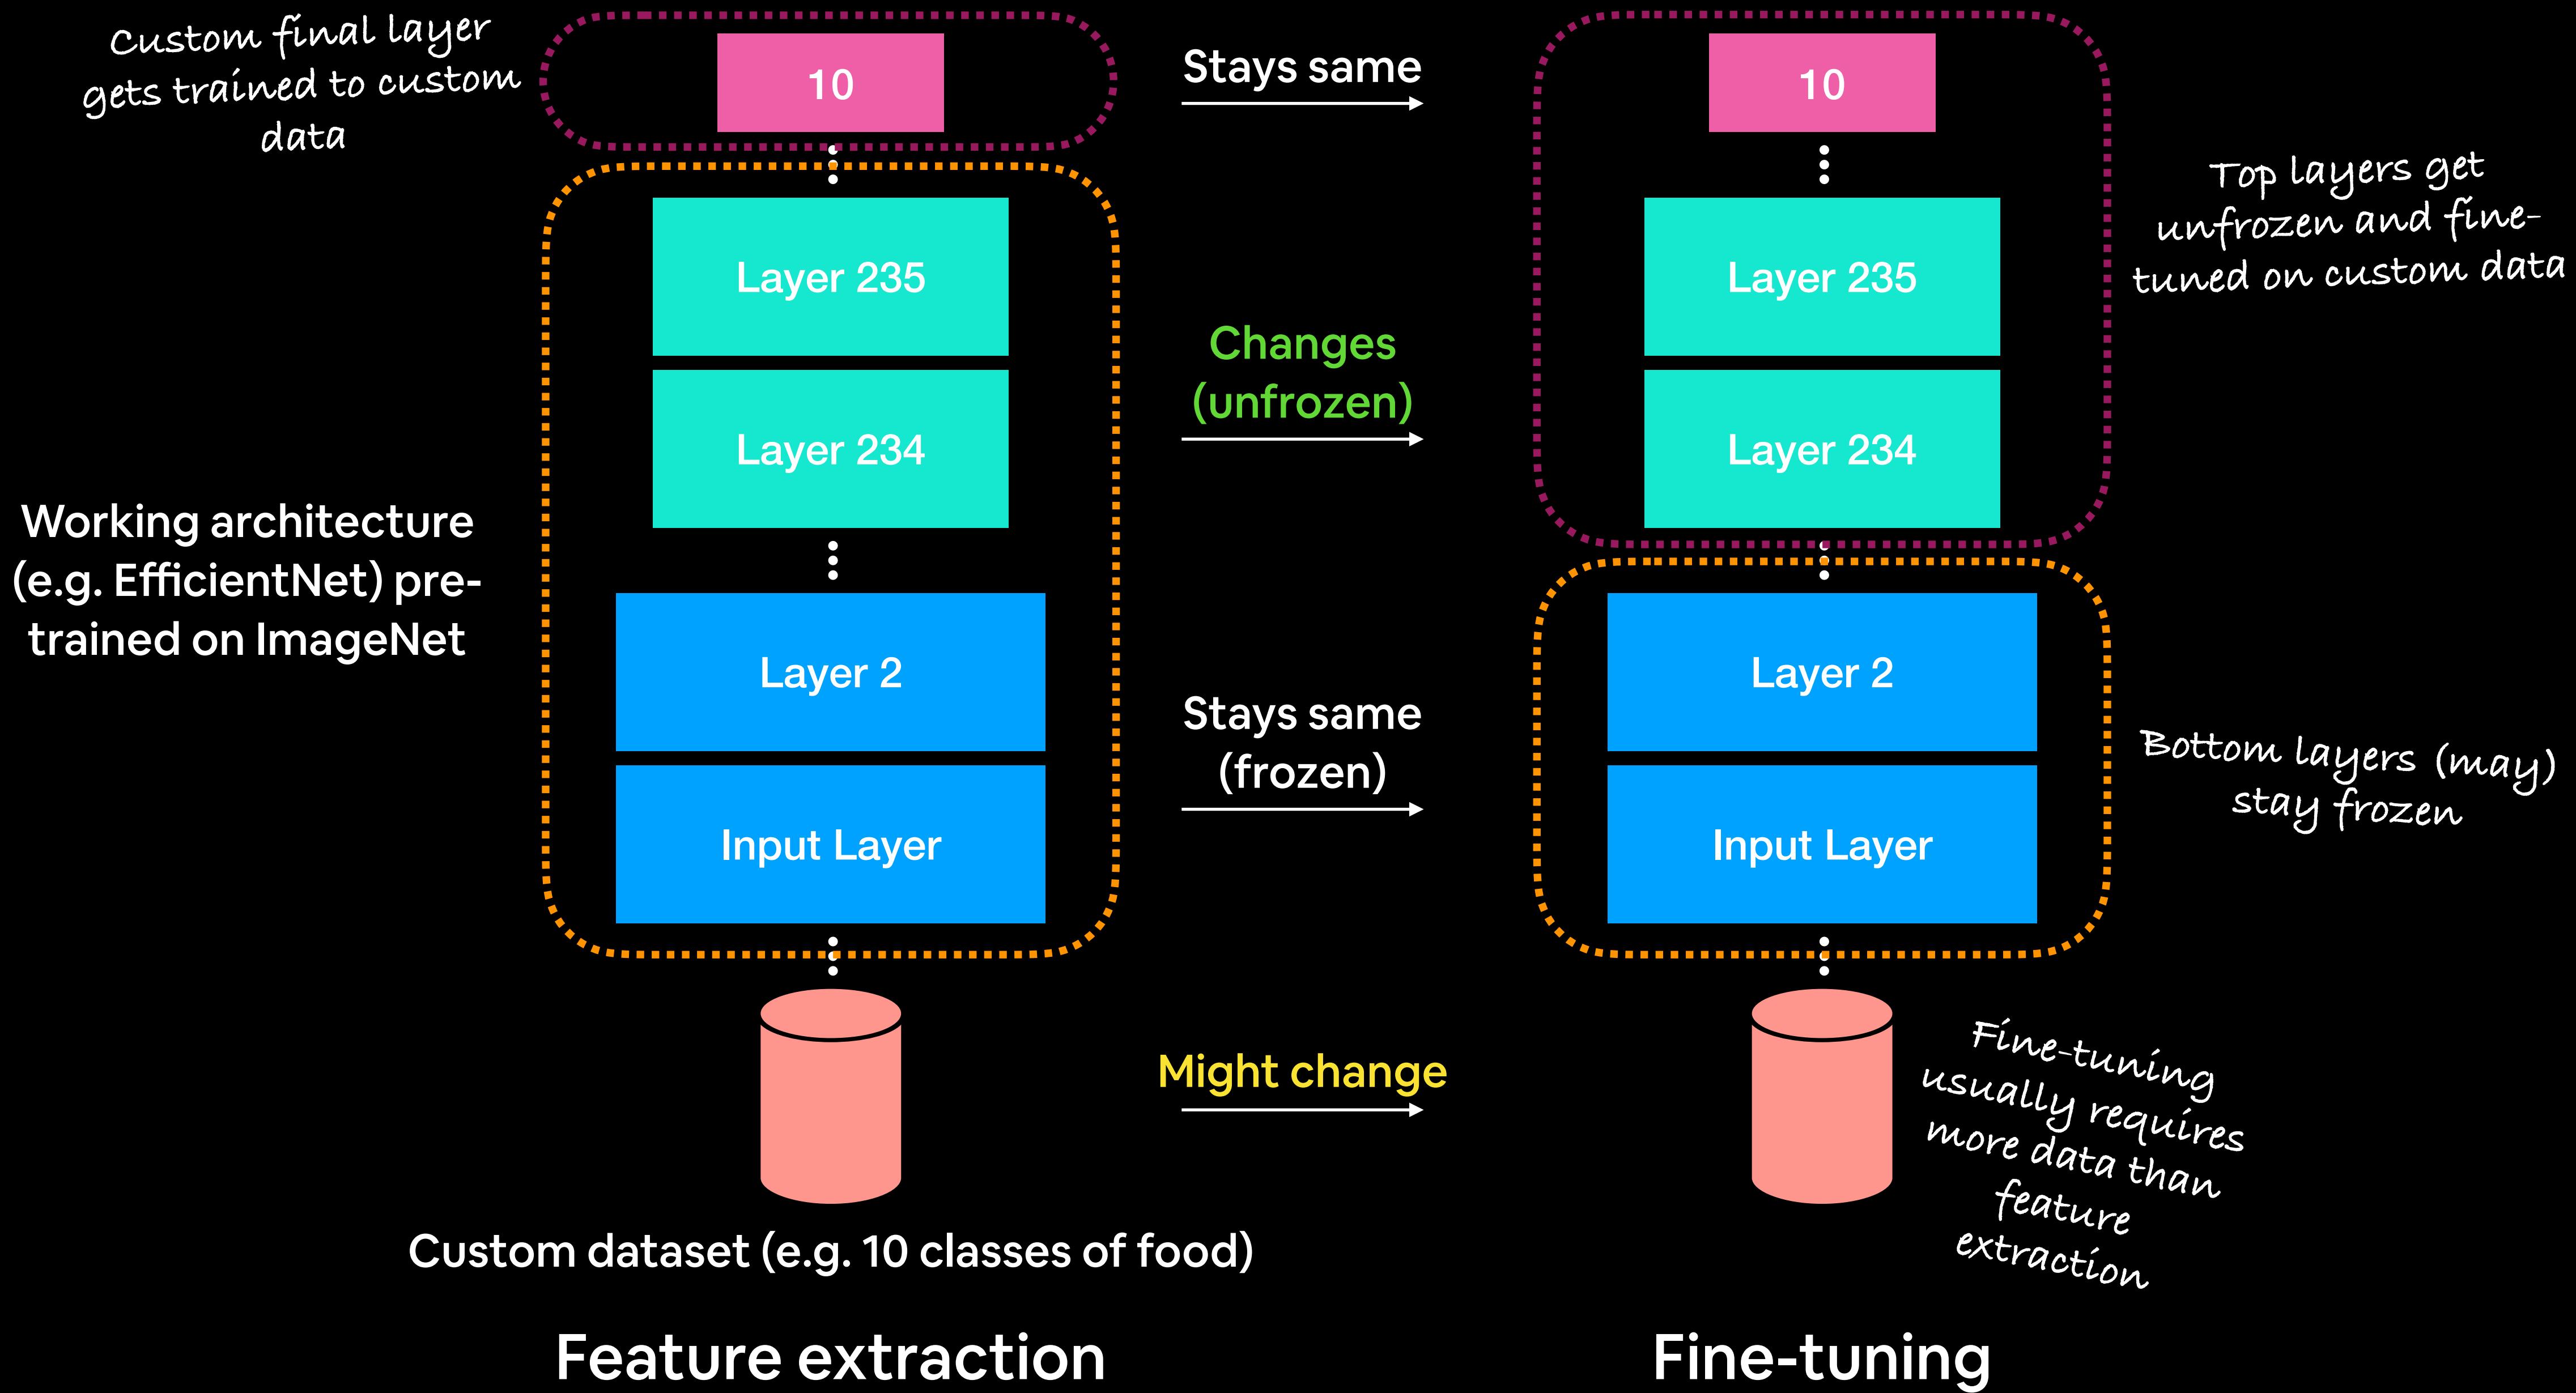
Learn patterns in a wide variety of images (using ImageNet)

EfficientNet architecture (already works really well on computer vision tasks)

Tune patterns/weights to our own problem (Food Vision)

Model performs better than from scratch

Feature extraction vs. Fine-tuning



What we're going to cover

(broadly)

- Introduce **fine-tuning transfer learning** with TensorFlow
- Introduce the **Keras Functional API** to build models
- Using a small dataset to experiment faster (e.g. 10% of training samples)
- **Data augmentation** (making your training set more diverse without adding samples)
- Running a series of experiments on our Food Vision data
- Introduce the **ModelCheckpoint callback** to save intermediate training results

(we'll be cooking up lots of code!)

How:



Let's code!

Dataset shapes

“Create batches of 32 images of size 224x224 split into red, green, blue colour channels.”

```
train_data_10_percent = tf.keras.preprocessing.image_dataset_from_directory(directory=train_dir,  
                                                               image_size=(224, 224),  
                                                               label_mode="categorical",  
                                                               batch_size=32)  
  
train_data_10_percent  
  
Found 750 files belonging to 10 classes.  
<BatchDataset shapes: ((None, 224, 224, 3), (None, 10)), types: (tf.float32, tf.float32)>
```

- Number of total samples (750 images, 75 per class)
- Number of classes (10 types of food)
- Batch size (default is 32)
- Image size (height, width)
- Number of colour channels (red, green, blue)
- Number of classes in label tensors (10 types of food)

Keras Sequential vs Functional API

Sequential API

```
# Creating a model with the Sequential API
sequential_model = tf.keras.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(64, activation="relu"),
    tf.keras.layers.Dense(64, activation="relu"),
    tf.keras.layers.Dense(10, activation="softmax")
], name="sequential_model")

sequential_model.compile(
    loss=tf.keras.losses.SparseCategoricalCrossentropy(),
    optimizer=tf.keras.optimizers.Adam(),
    metrics=[ "accuracy" ]
)

sequential_model.fit(x_train, y_train,
                     batch_size=32,
                     epochs=5)
```

Functional API

```
# Creating a model with the Functional API
inputs = tf.keras.layers.Input(shape=(28, 28))
x = tf.keras.layers.Flatten()(inputs)
x = tf.keras.layers.Dense(64, activation="relu")(x)
x = tf.keras.layers.Dense(64, activation="relu")(x)
outputs = tf.keras.layers.Dense(10, activation="softmax")(x)
functional_model = tf.keras.Model(inputs, outputs, name="functional_model")

functional_model.compile(
    loss=tf.keras.losses.SparseCategoricalCrossentropy(),
    optimizer=tf.keras.optimizers.Adam(),
    metrics=[ "accuracy" ]
)

functional_model.fit(x_train, y_train,
                     batch_size=32,
                     epochs=5)
```

compiling and fitting stays the same

- Similarities: **compiling, fitting, evaluating**
- Differences: **model construction** (the **Functional API is more flexible** and able to produce more sophisticated models)

Building a feature extraction model with the Keras Functional API

```
# 1. Create base model with tf.keras.applications
base_model = tf.keras.applications.EfficientNetB0(include_top=False)

# 2. Freeze the base model (so the pre-learned patterns remain)
base_model.trainable = False

# 3. Create inputs into the base model
inputs = tf.keras.layers.Input(shape=(224, 224, 3), name="input_layer")

# 4. If using ResNet50V2, add this to speed up convergence, remove for EfficientNet
# x = tf.keras.layers.experimental.preprocessing.Rescaling(1./255)(inputs)

# 5. Pass the inputs to the base_model (note: using tf.keras.applications, EfficientNet inputs don't have to be normalized)
x = base_model(inputs)
# Check data shape after passing it to base_model
print(f"Shape after base_model: {x.shape}")

# 6. Average pool the outputs of the base model (aggregate all the most important information)
x = tf.keras.layers.GlobalAveragePooling2D(name="global_average_pooling_layer")(x)
print(f"After GlobalAveragePooling2D(): {x.shape}")

# 7. Create the output activation layer
outputs = tf.keras.layers.Dense(10, activation="softmax", name="output_layer")(x)

# 8. Combine the inputs with the outputs into a model
model_0 = tf.keras.Model(inputs, outputs)

# 9. Compile the model
model_0.compile(loss='categorical_crossentropy',
                 optimizer=tf.keras.optimizers.Adam(),
                 metrics=[ "accuracy"])

# 10. Fit the model (we use less steps for validation so it's faster)
history_10_percent = model_0.fit(train_data_10_percent,
                                   epochs=5,
                                   steps_per_epoch=len(train_data_10_percent),
                                   validation_data=test_data_10_percent,
                                   # Go through less of the validation data so epochs are faster (we want faster experiments!)
                                   validation_steps=int(0.25 * len(test_data_10_percent)),
                                   # Track our model's training logs for visualization later
                                   callbacks=[create_tensorboard_callback("transfer_learning", "10_percent_feature_extract")])
```

Useful computer vision architectures

- `tf.keras.applications` and `keras.applications` have many of the **most popular and best performing computer vision architectures built-in & pre-trained**, ready to use for your own problems

The screenshot shows the Keras Applications page. The left sidebar has a red background for the 'Keras Applications' section. The main content area includes a search bar, a breadcrumb trail ('» Keras API reference / Keras Applications'), and a table of available models.

Keras Applications

Keras Applications are deep learning models that are made available alongside pre-trained weights. These models can be used for prediction, feature extraction, and fine-tuning.

Weights are downloaded automatically when instantiating a model. They are stored at `~/.keras/models/`.

Upon instantiation, the models will be built according to the image data format set in your Keras configuration file at `~/.keras/keras.json`. For instance, if you have set `image_data_format=channels_last`, then any model loaded from this repository will get built according to the TensorFlow data format convention, "Height-Width-Depth".

Available models

Model	Size	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth
Xception	88 MB	0.790	0.945	22,910,480	126
VGG16	528 MB	0.713	0.901	138,357,544	23
VGG19	549 MB	0.713	0.900	143,667,240	26
ResNet50	98 MB	0.749	0.921	25,636,712	-
ResNet101	171 MB	0.764	0.928	44,707,176	-
ResNet152	232 MB	0.766	0.931	60,419,944	-
ResNet50V2	98 MB	0.760	0.930	25,613,800	-

Source: <https://keras.io/api/applications/>

The screenshot shows the tf.keras.applications module page. The top navigation bar includes the TensorFlow logo and language selection (English). The main content area includes a table of contents, a sidebar with 'TensorFlow 1 version' button, and a note about Keras Applications.

Module: tf.keras.applications

Table of contents

Modules

Functions

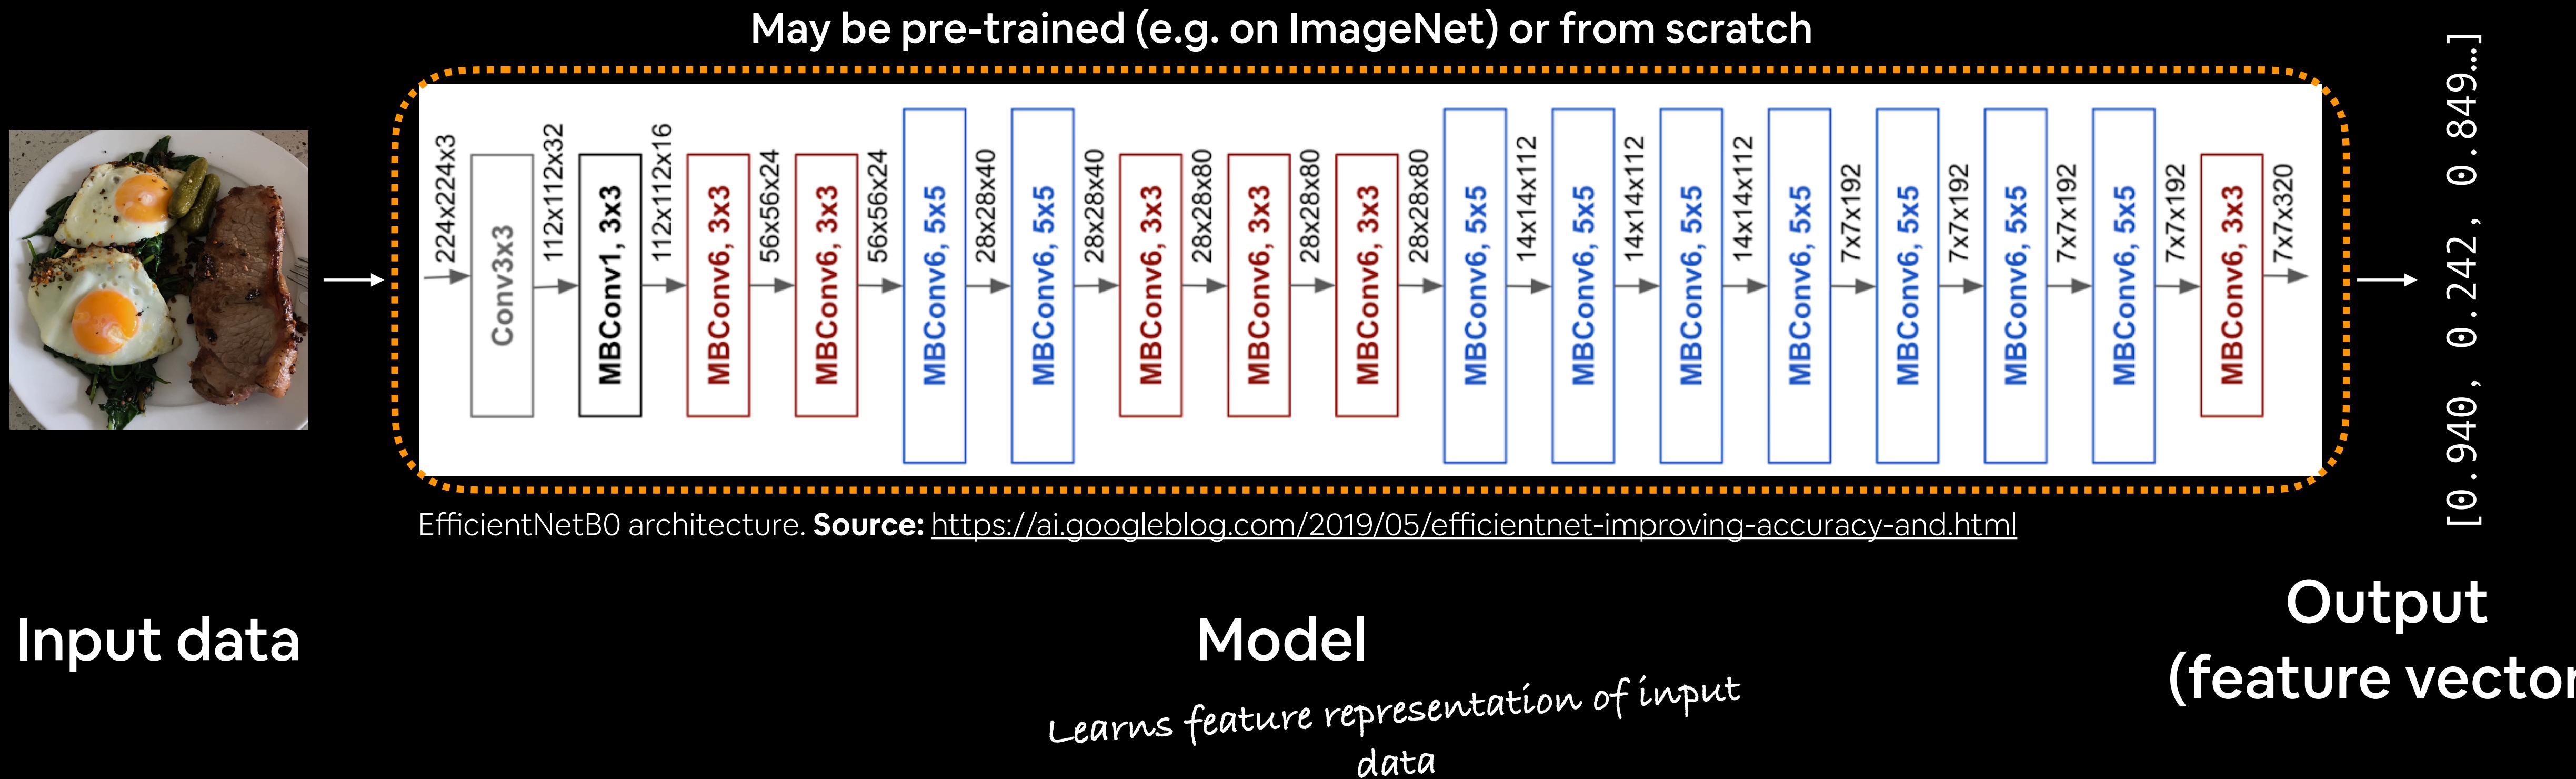
TensorFlow 1 version

Keras Applications are canned architectures with pre-trained weights.

Source: https://www.tensorflow.org/api_docs/python/tf/keras/applications

What is a feature vector?

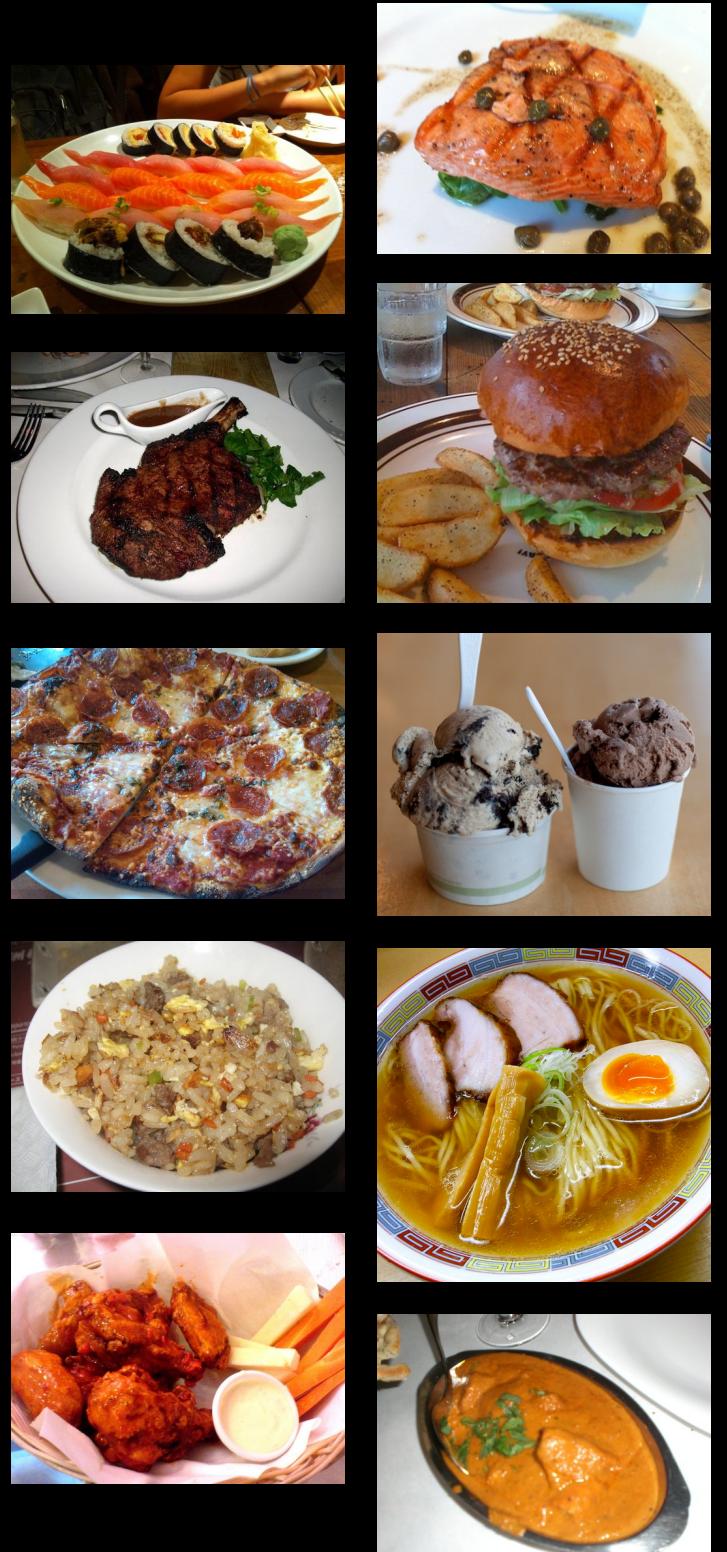
- A feature vector is **a learned representation of the input data** (a compressed form of the input data based on how the model sees it)



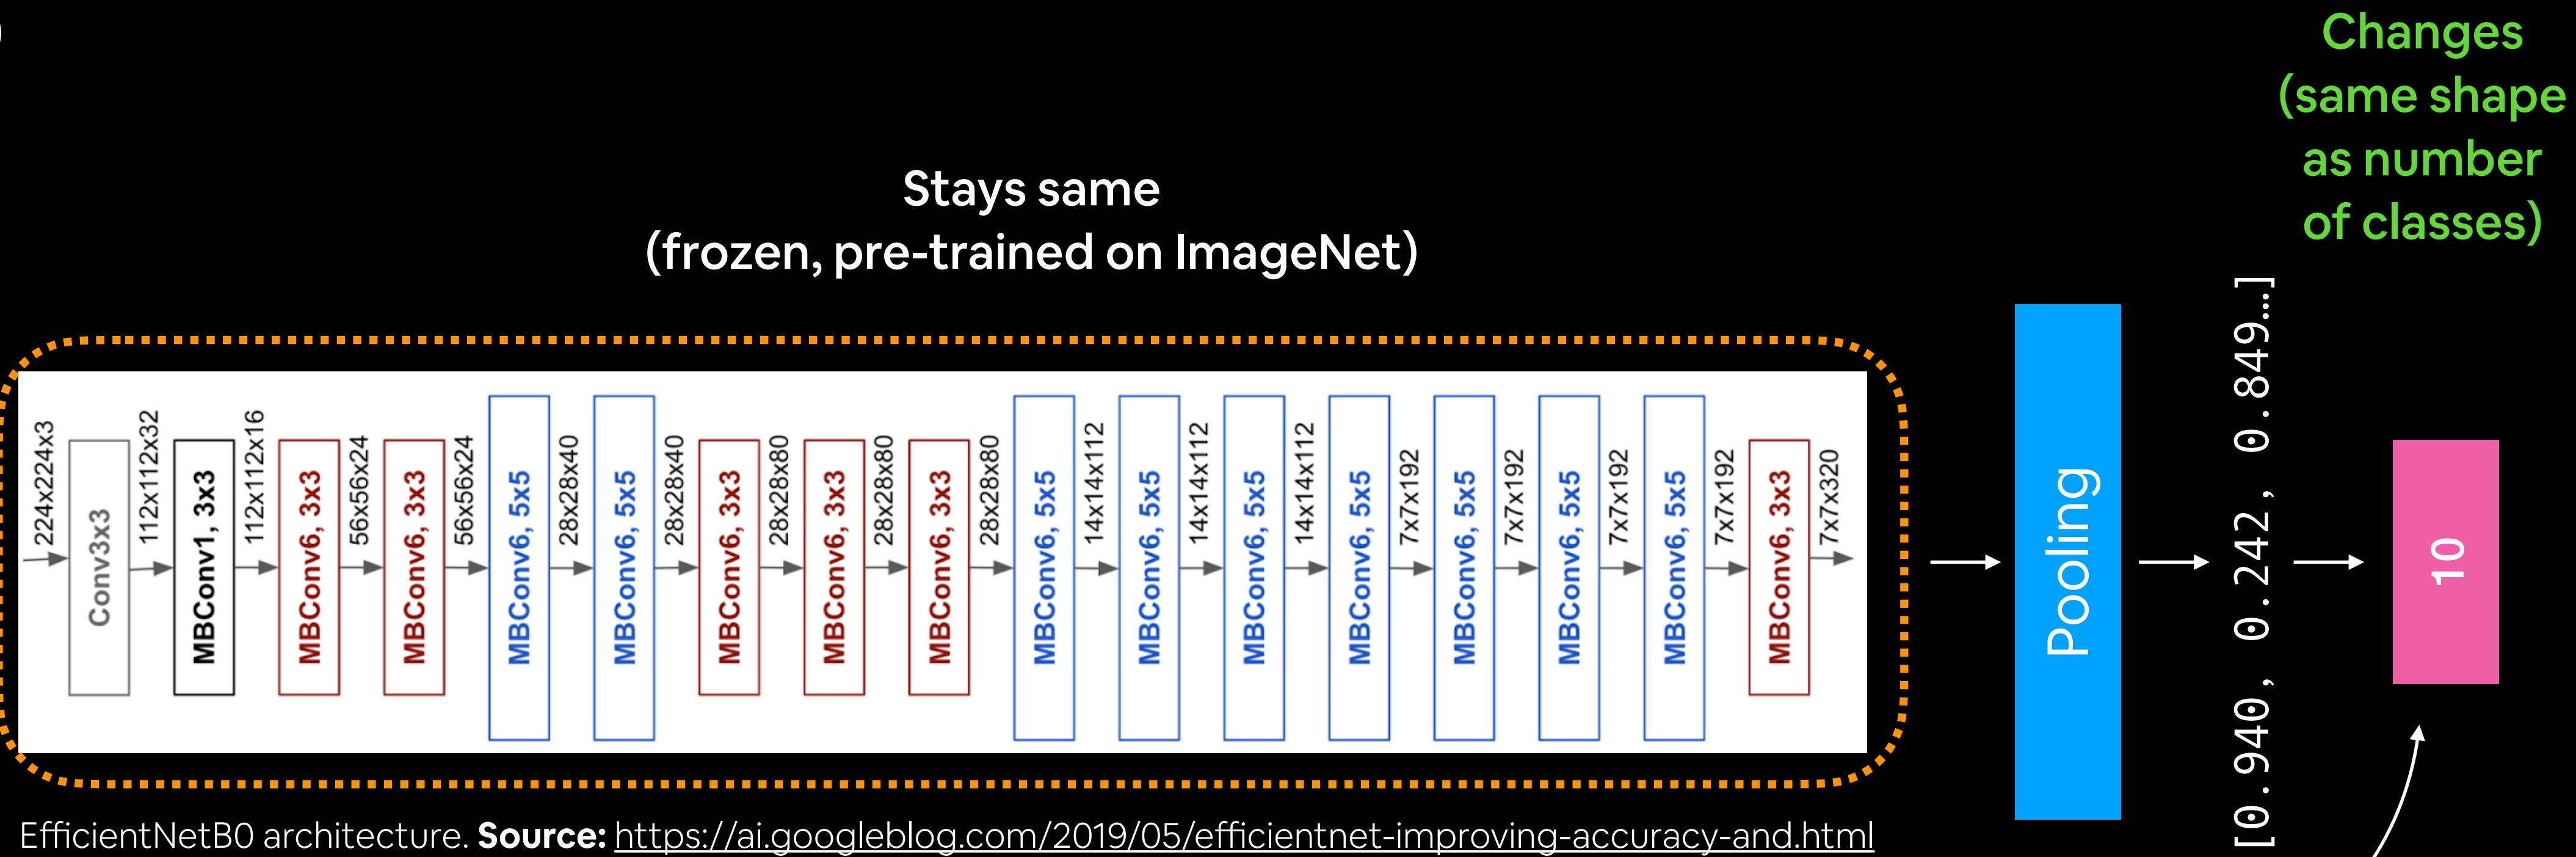
EfficientNet feature extractor

Input data

(10 classes of Food101)



Stays same
(frozen, pre-trained on ImageNet)

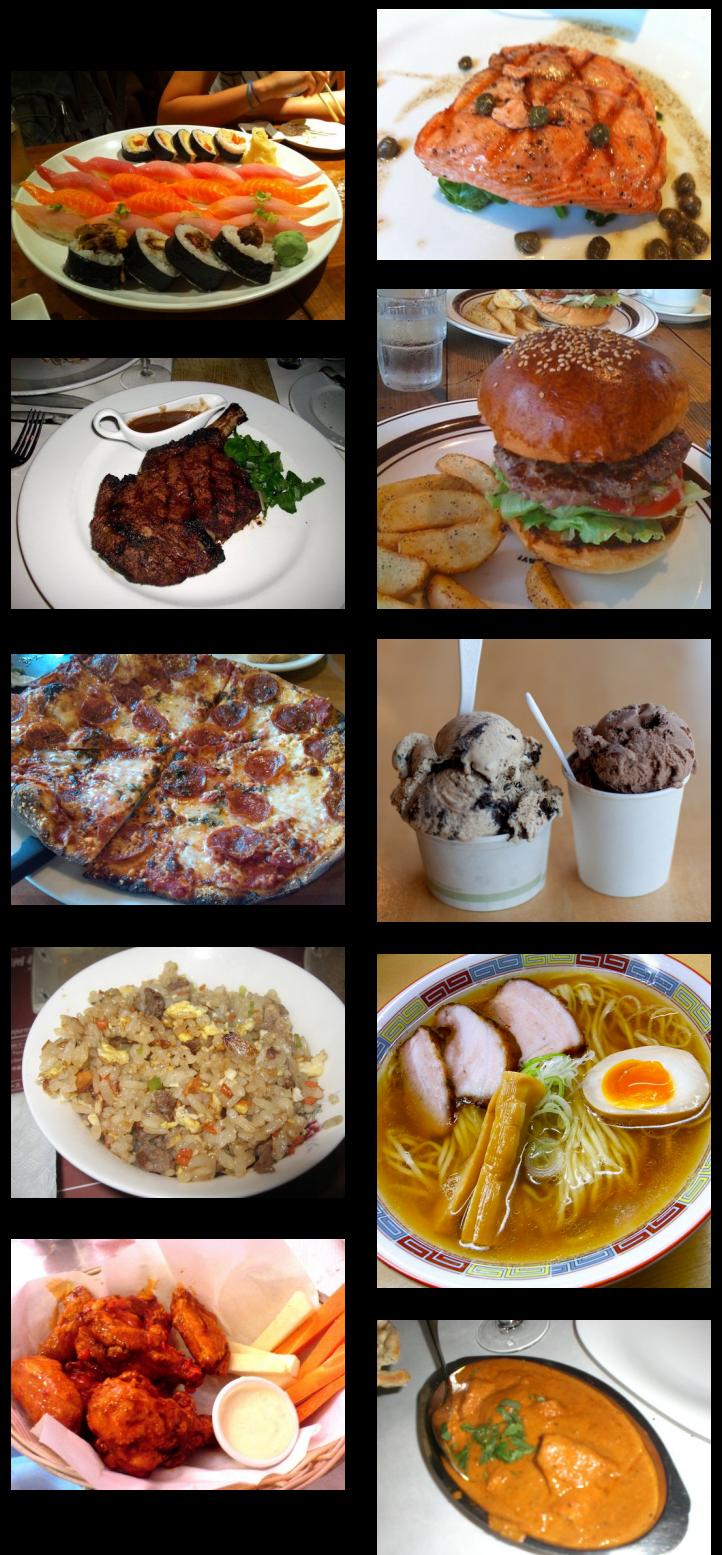


EfficientNetB0 architecture. **Source:** <https://ai.googleblog.com/2019/05/efficientnet-improving-accuracy-and.html>

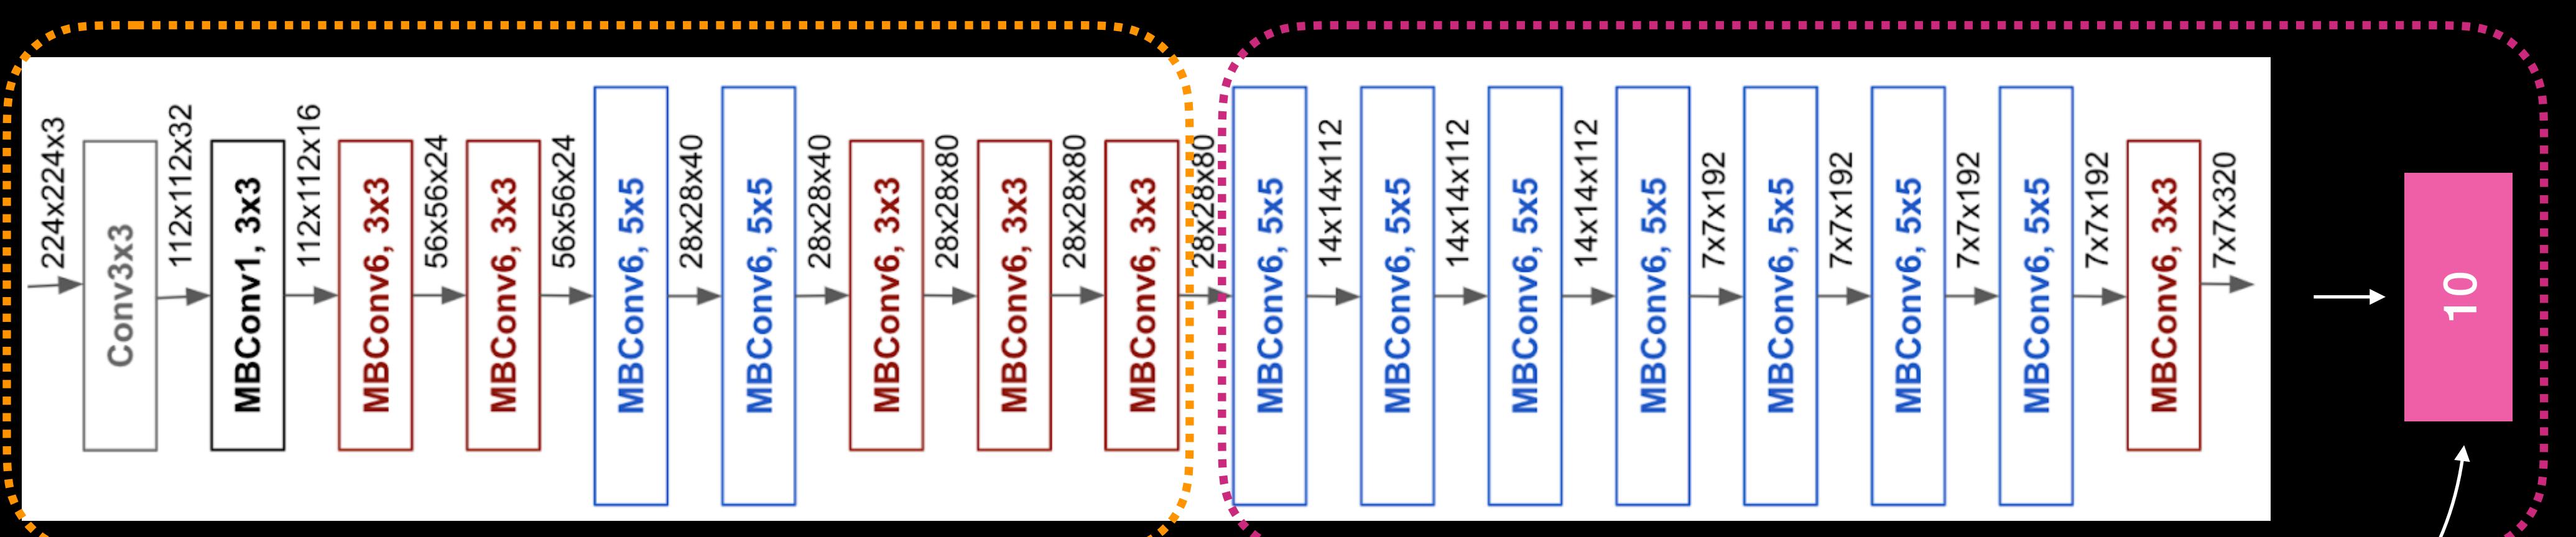
Fully-connected
(dense) classifier layer

EfficientNet fine-tuning

Input data
(10 classes of Food101)



Stays same
(frozen, pre-trained on ImageNet)



EfficientNetB0 architecture. **Source:** <https://ai.googleblog.com/2019/05/efficientnet-improving-accuracy-and.html>

Bottom layers tend to stay frozen (or are last to get unfrozen)

Layers closer to the output layer get unfrozen/fine-tuned first

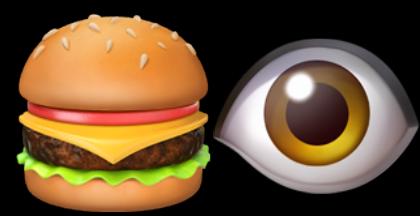
Changes
(unfrozen, gets fine-tuned on custom data)

Fully-connected
(dense) classifier layer

10

Modelling experiments we're running

Experiment	Data	Preprocessing	Model
Model 0 (baseline)	10 classes of Food101 data (random 10% training data only)	None	Feature Extractor: EfficientNetB0 (pre-trained on ImageNet, all layers frozen) with no top
Model 1	10 classes of Food101 data (random 1% training data only)	Random Flip, Rotation, Zoom, Height, Width data augmentation	Same as Model 0
Model 2	Same as Model 0	Same as Model 1	Same as Model 0
Model 3	Same as Model 0	Same as Model 1	Fine-tuning: Model 2 (EfficientNetB0 pre-trained on ImageNet) with top layer trained on custom data, top 10 layers unfrozen
Model 4	10 classes of Food101 data (100% training data)	Same as Model 1	Same as Model 3

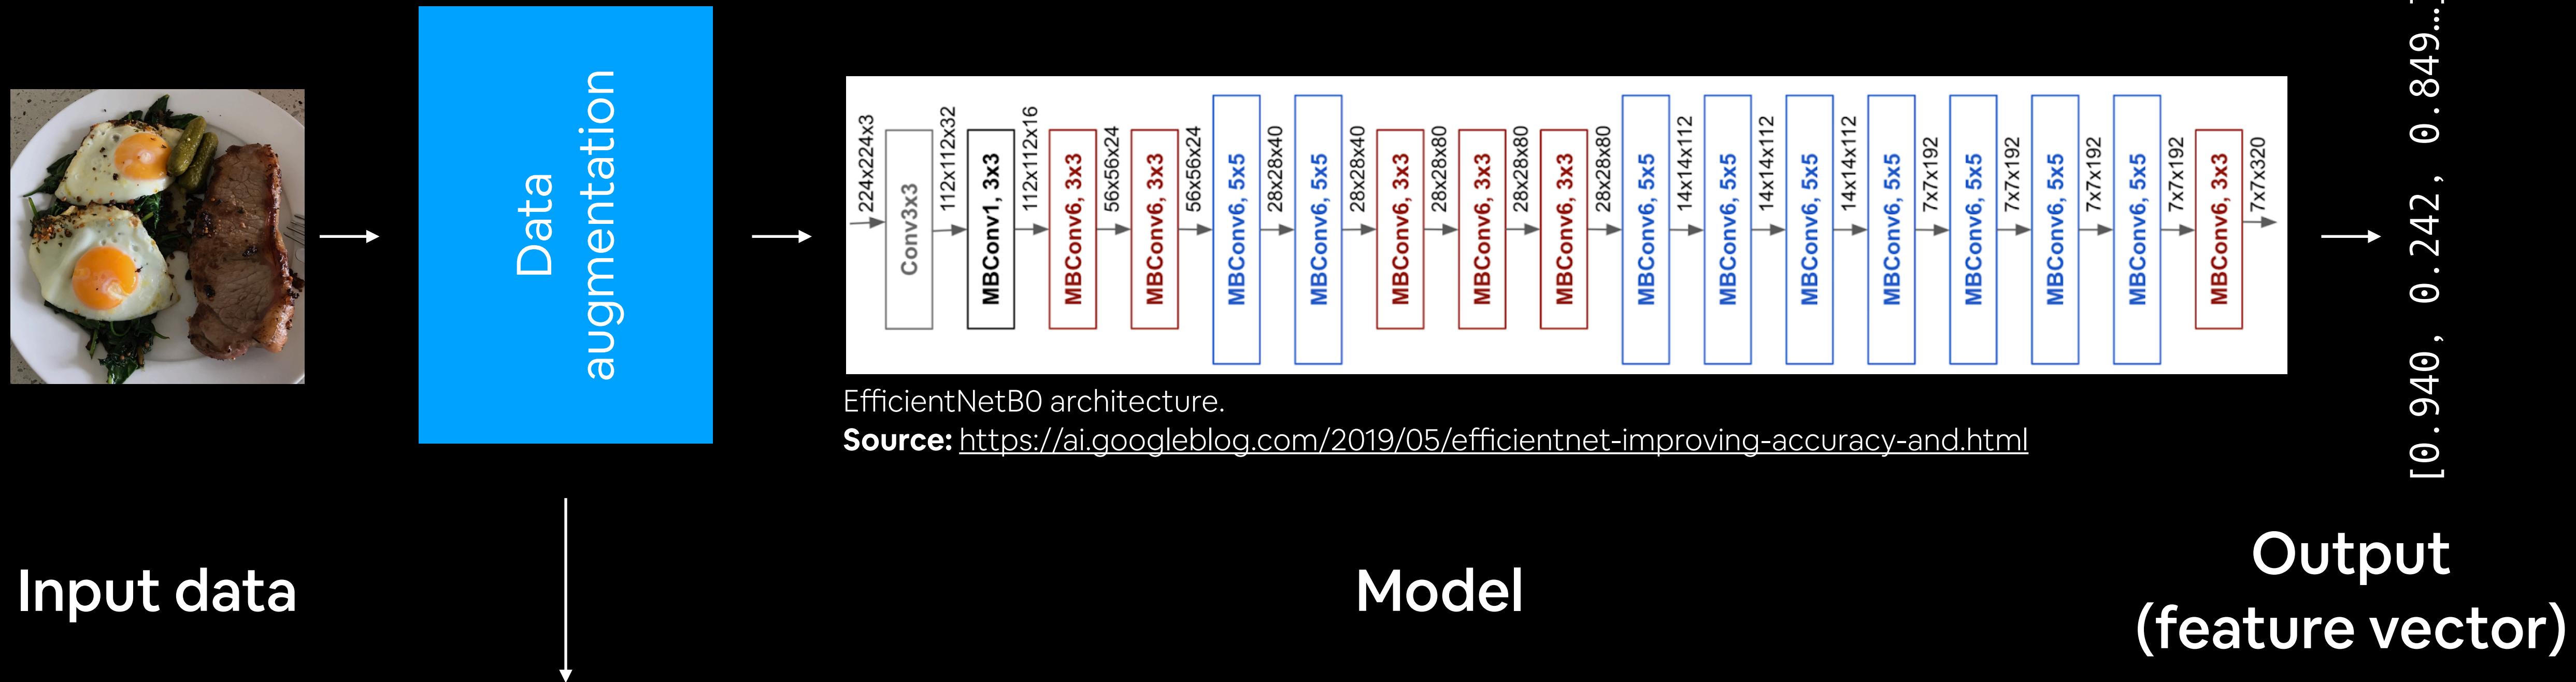


Food Vision: Dataset(s) we're using

Note: For randomly selected data, the Food101 dataset was downloaded and modified using the [Image Data Modification Notebook](#)

Dataset Name	Source	Classes	Training data	Testing data
pizza_steak	Food101	Pizza, steak (2)	750 images of pizza and steak (same as original Food101 dataset)	250 images of pizza and steak (same as original Food101 dataset)
10_food_classes_1_percent	Same as above	Chicken curry, chicken wings, fried rice, grilled salmon, hamburger, ice cream, pizza, ramen, steak, sushi (10)	7 randomly selected images of each class (1% of original training data)	250 images of each class (same as original Food101 dataset)
10_food_classes_10_percent	Same as above	Same as above	75 randomly selected images of each class (10% of original training data)	Same as above
10_food_classes_100_percent	Same as above	Same as above	750 images of each class (100% of original training data)	Same as above
101_food_classes_10_percent	Same as above	All classes from Food101 (101)	75 images of each class (10% of original Food101 dataset)	250 images of each class (same as original Food101 dataset)

Data augmentation as a layer



```
# Build data augmentation layer
data_augmentation = tf.keras.Sequential([
    preprocessing.RandomFlip('horizontal'),
    preprocessing.RandomHeight(0.2),
    preprocessing.RandomWidth(0.2),
    preprocessing.RandomZoom(0.2),
    preprocessing.RandomRotation(0.2),
], name="data_augmentation")
```

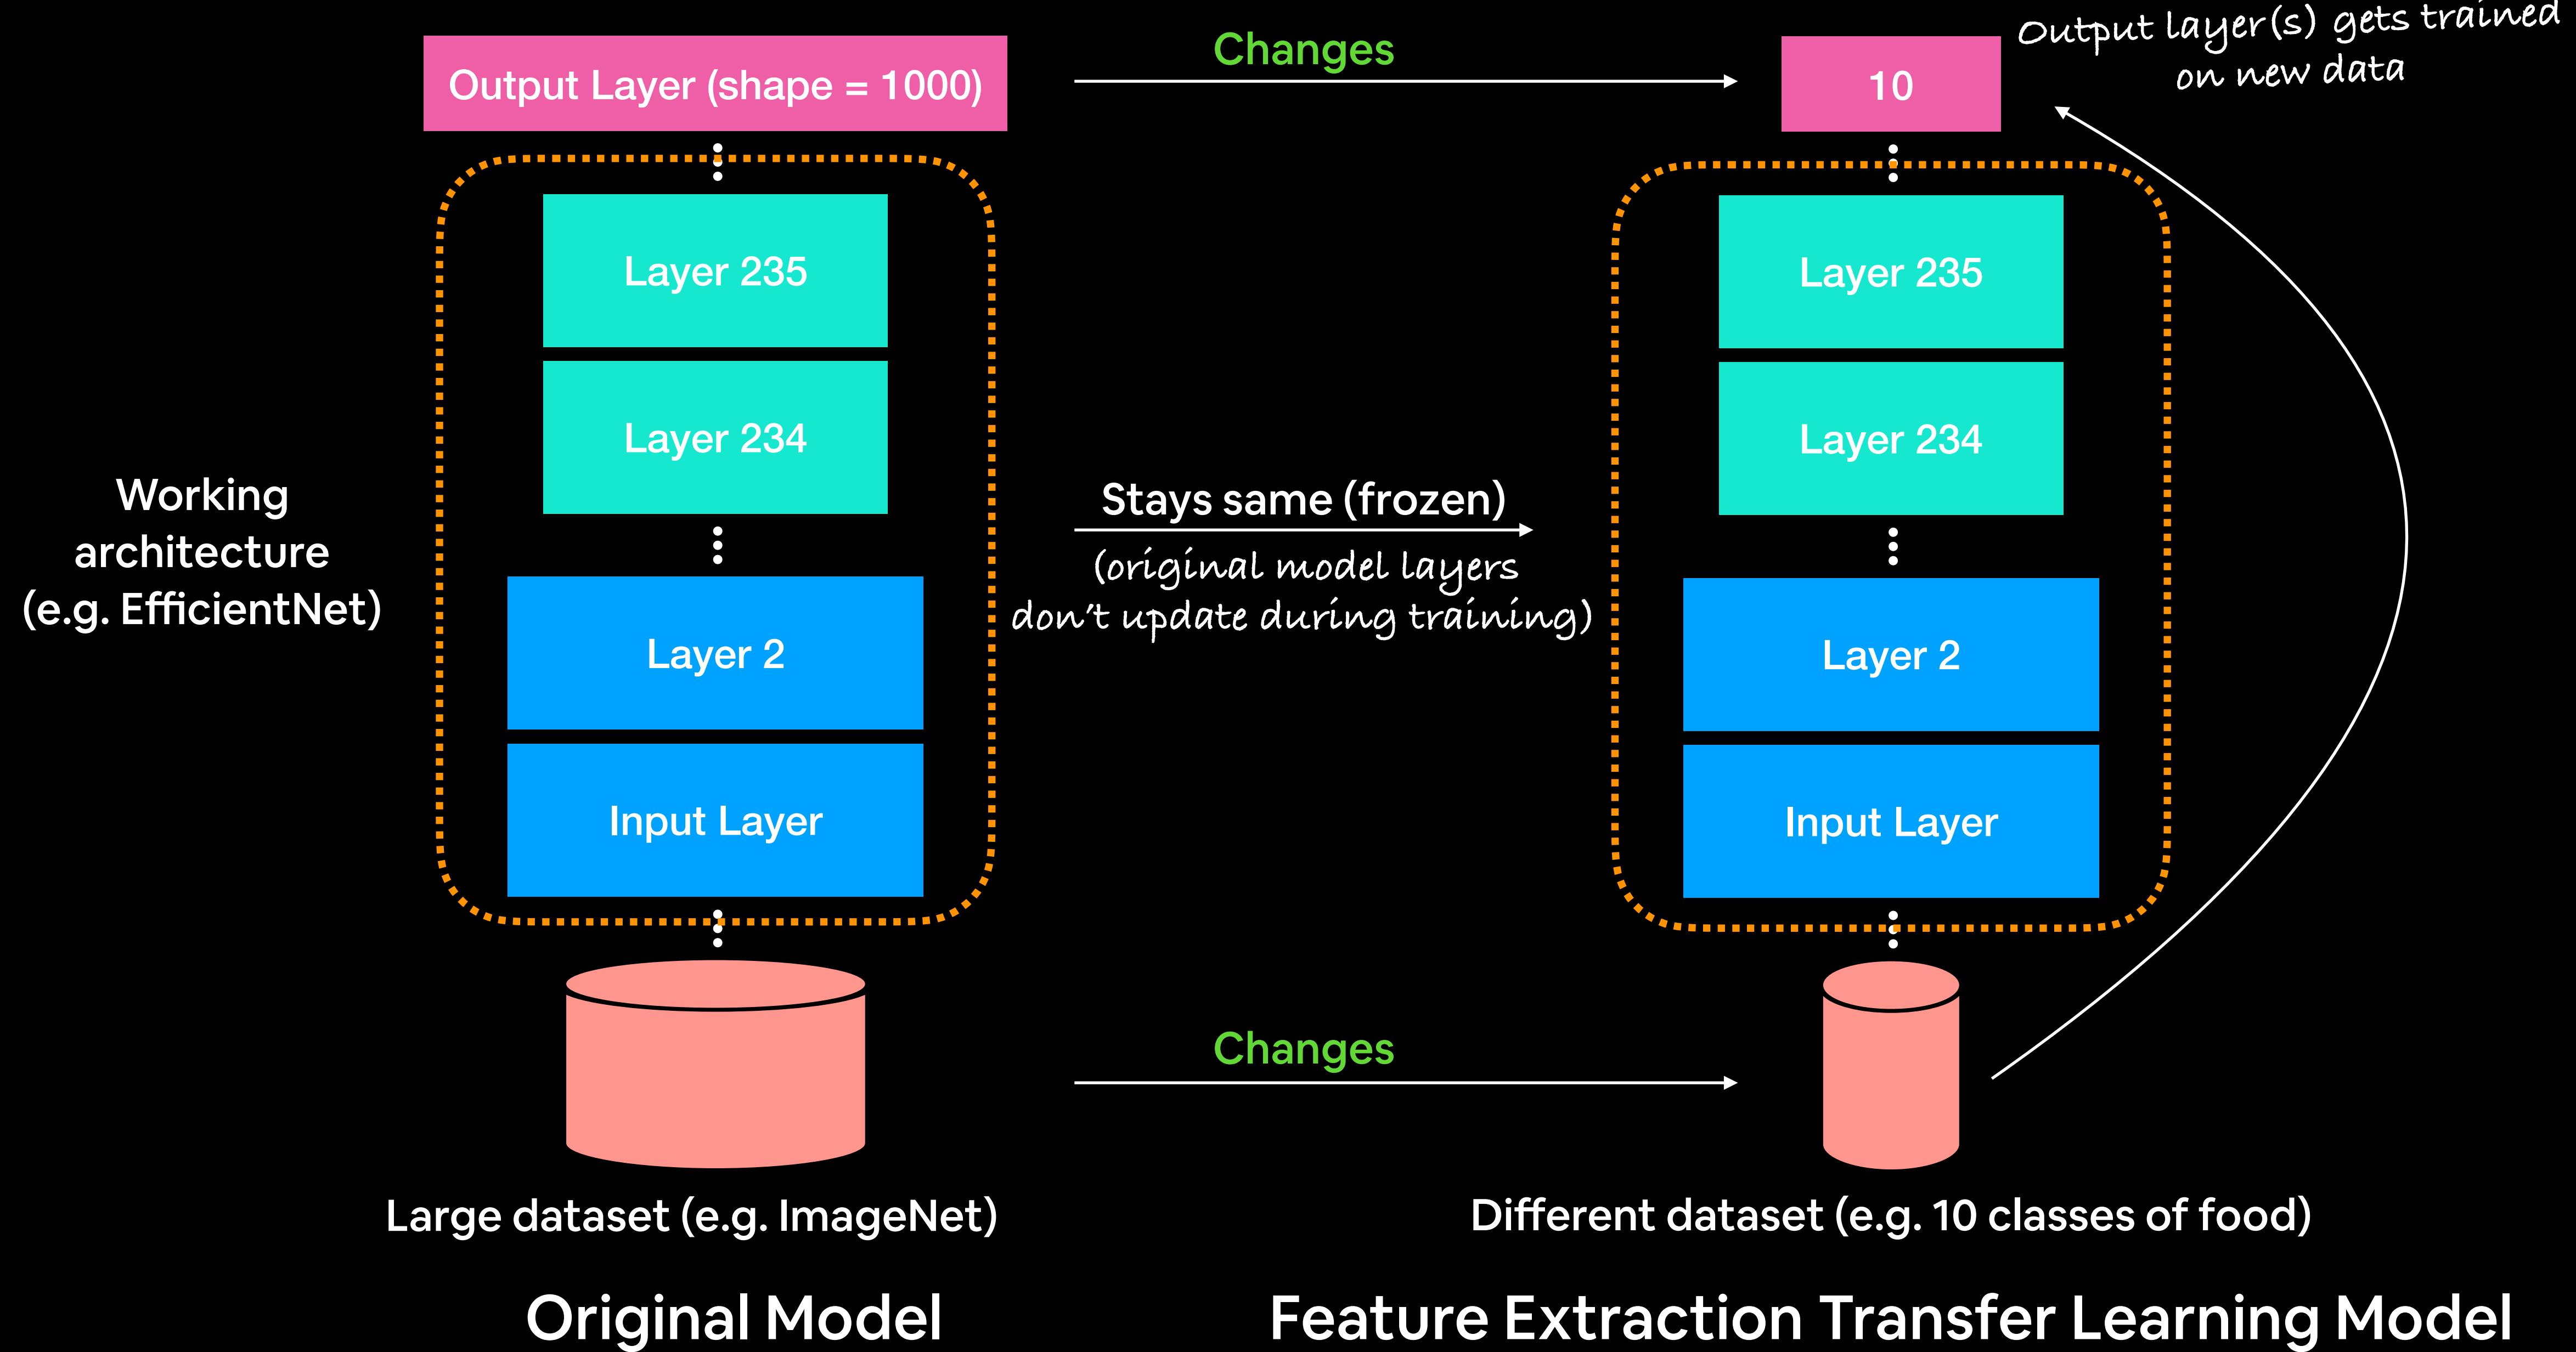
When passed as a layer to a model data augmentation is automatically **turned on during training** (augments training data) but **turned off during inference** (does not augment testing data or new unseen data).

What are callbacks?

- Callbacks are a tool which can **add helpful functionality** to your models during training, evaluation or inference
- Some popular callbacks include:

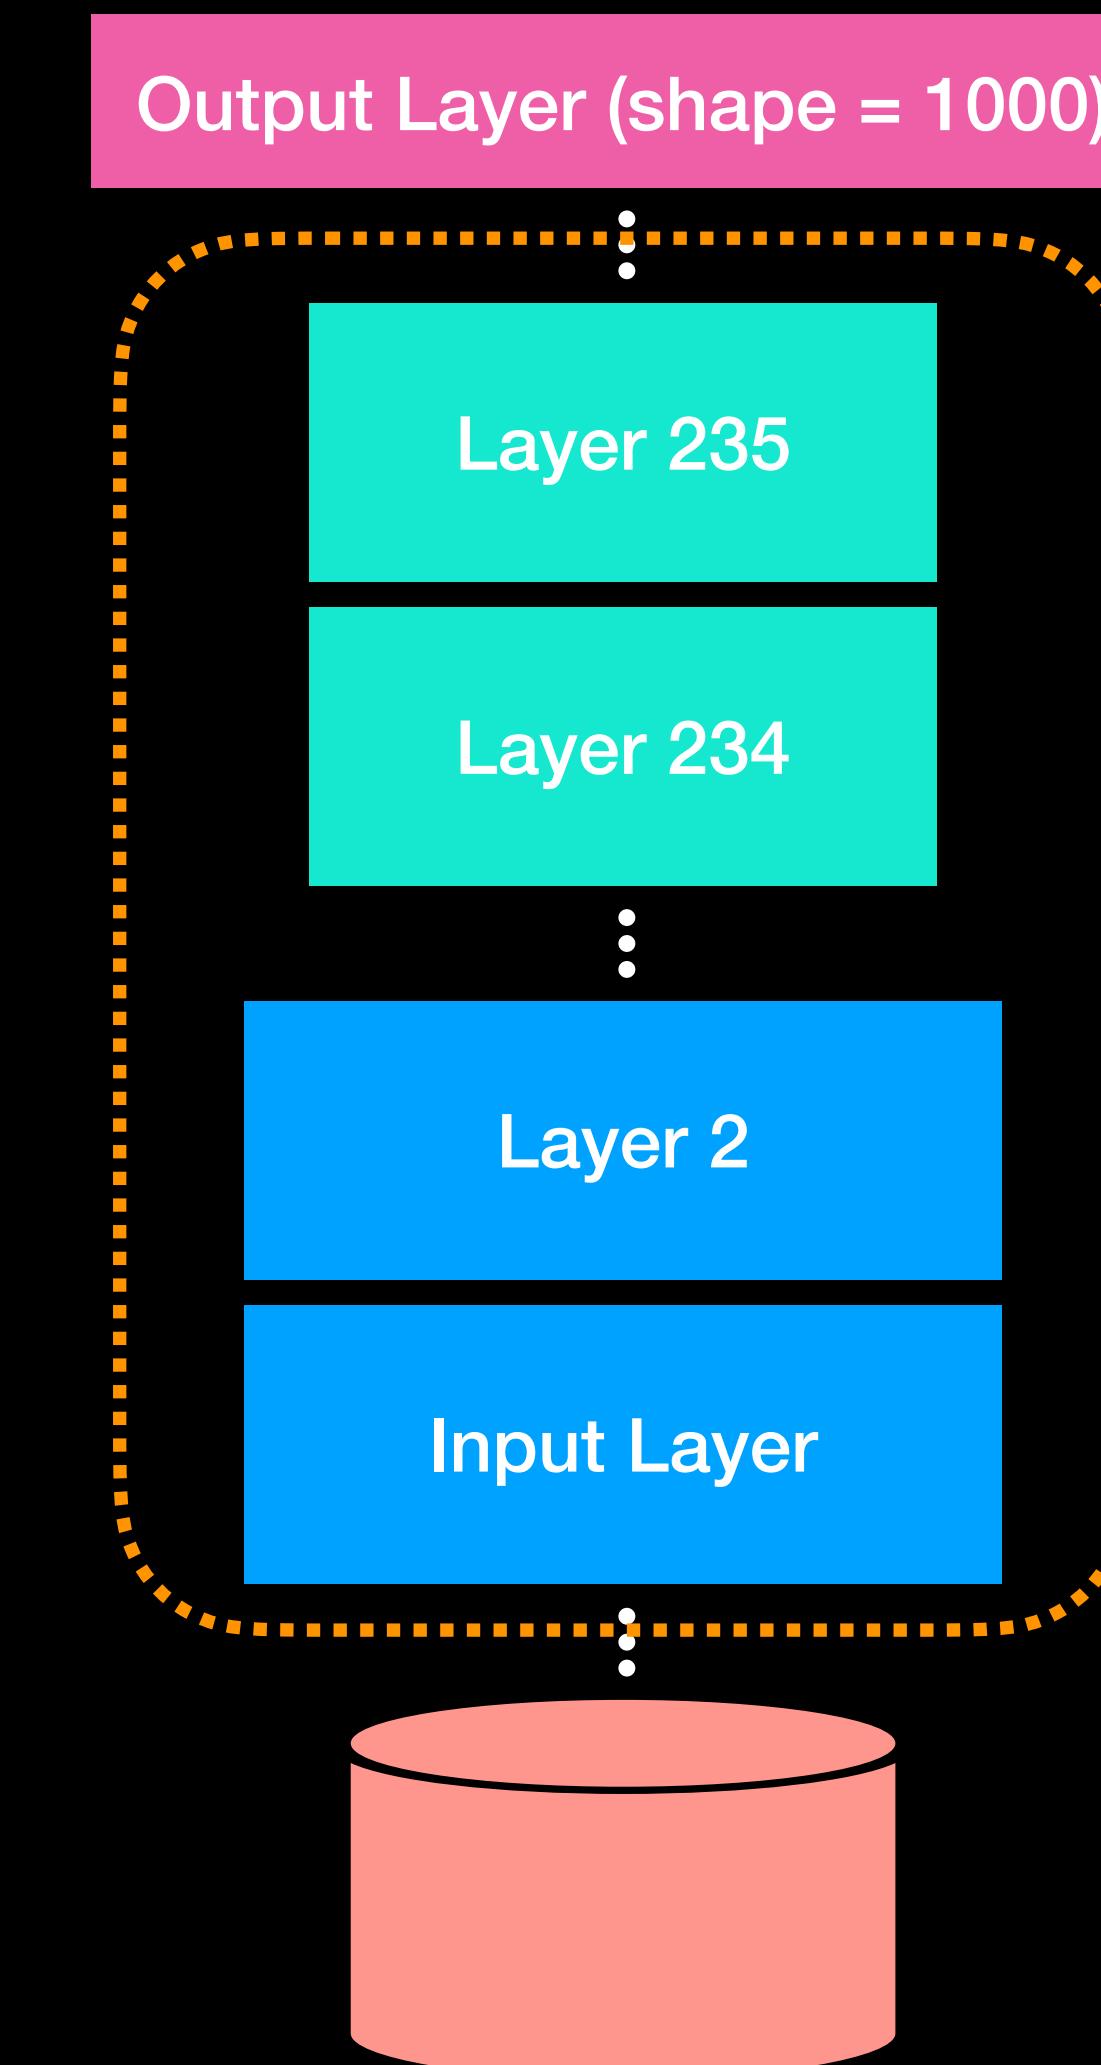
Callback name	Use case	Code
<u>TensorBoard</u>	Log the performance of multiple models and then view and compare these models in a visual way on TensorBoard (a dashboard for inspecting neural network parameters). Helpful to compare the results of different models on your data.	<code>tf.keras.callbacks.TensorBoard()</code>
<u>Model checkpointing</u>	Save your model as it trains so you can stop training if needed and come back to continue off where you left. Helpful if training takes a long time and can't be done in one sitting.	<code>tf.keras.callbacks.ModelCheckpoint()</code>
<u>Early stopping</u>	Leave your model training for an arbitrary amount of time and have it stop training automatically when it ceases to improve. Helpful when you've got a large dataset and don't know how long training will take.	<code>tf.keras.callbacks.EarlyStopping()</code>

Original Model vs. Feature Extraction



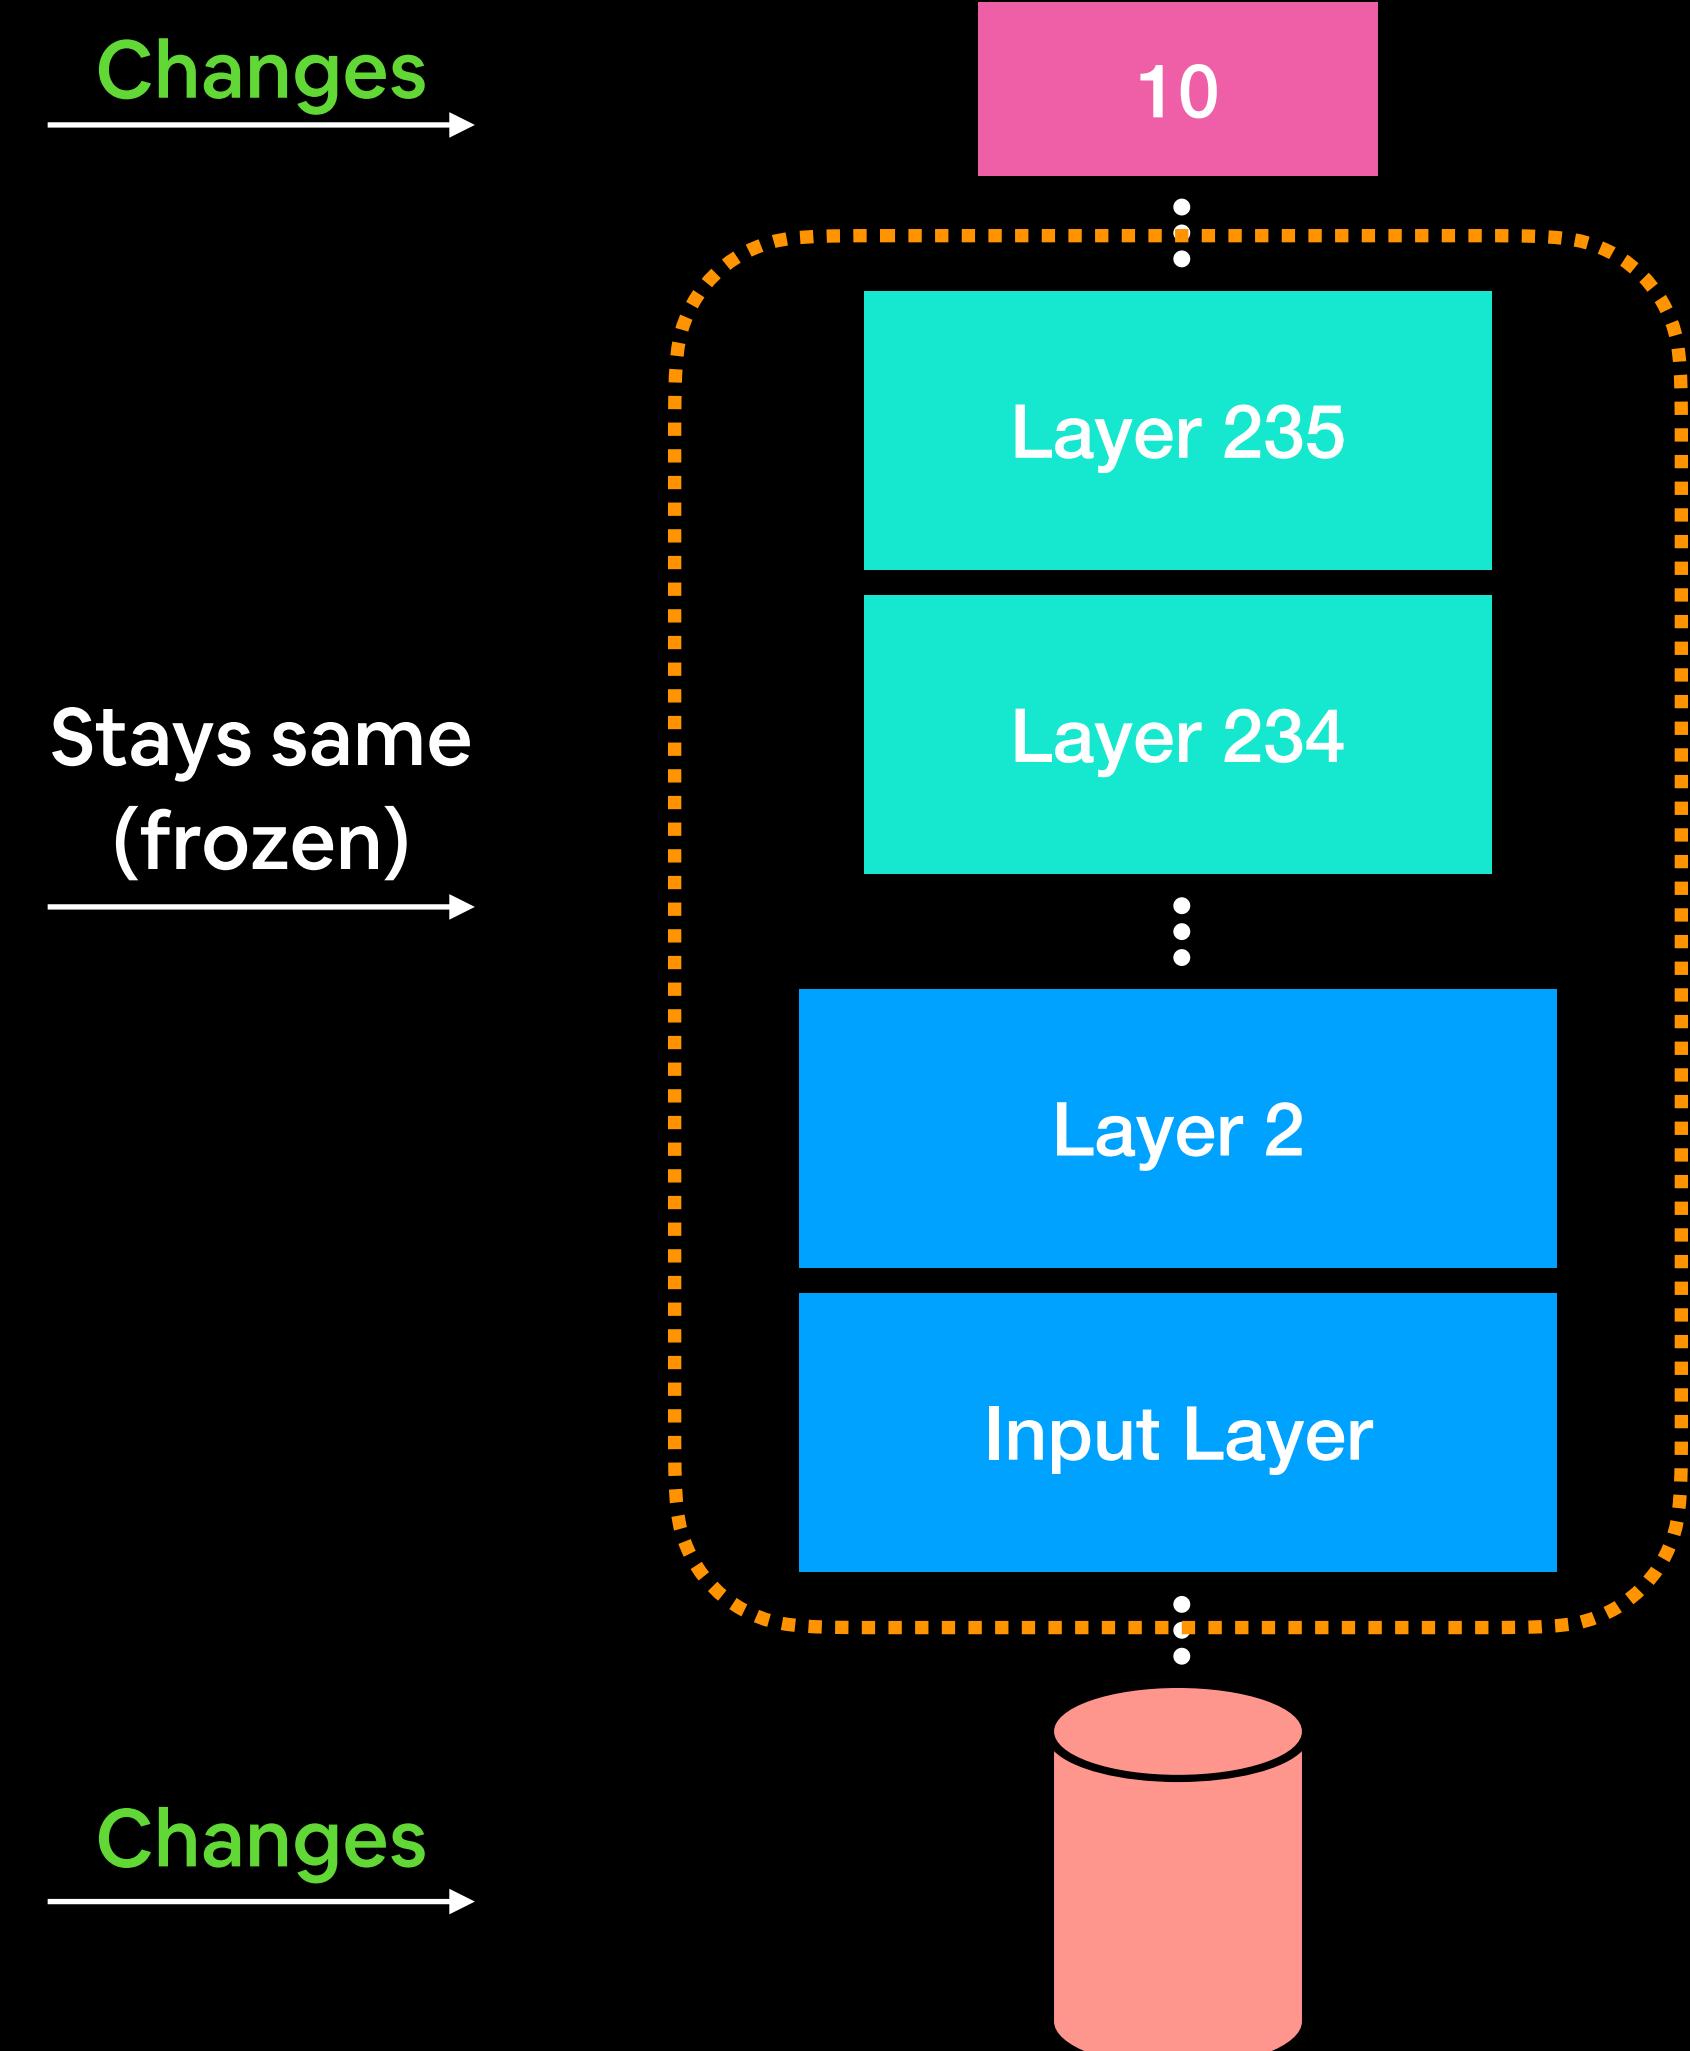
Kinds of Transfer Learning

Top layers get trained
on new data



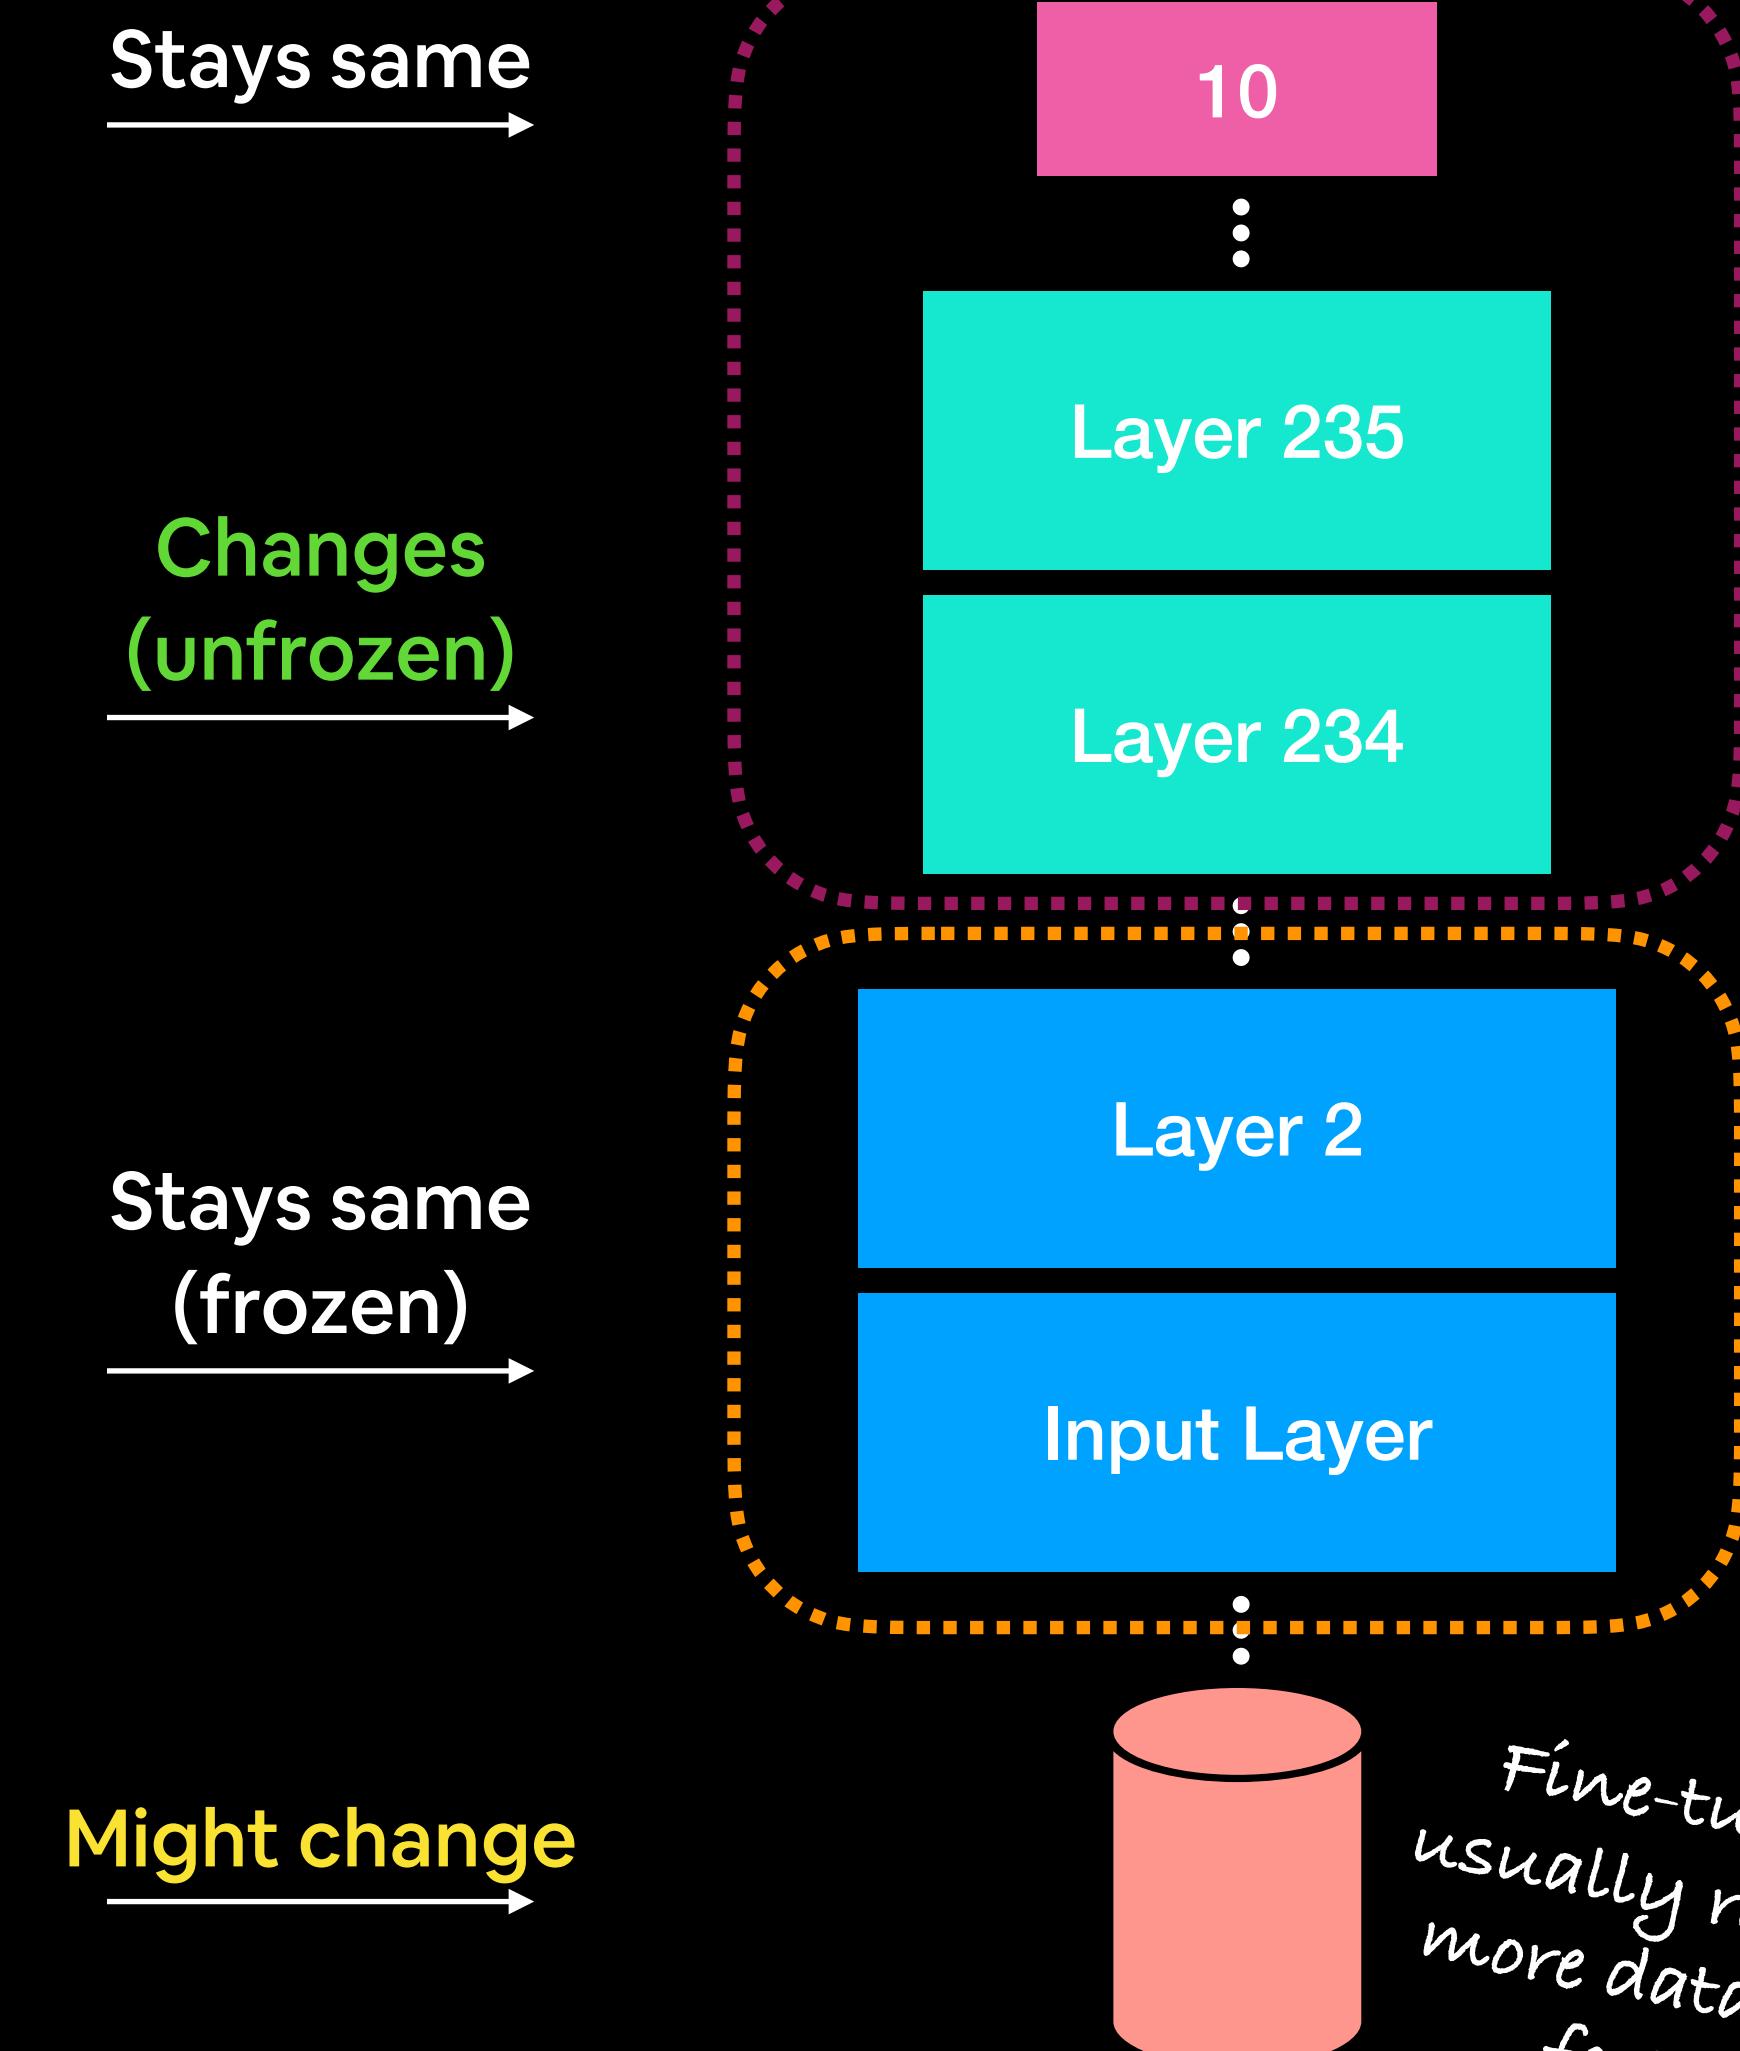
Large dataset (e.g. ImageNet)

Original Model



Different dataset (e.g. 10 classes of food)

Feature Extraction



Might change

Fine-tuning
usually requires
more data than
feature
extraction

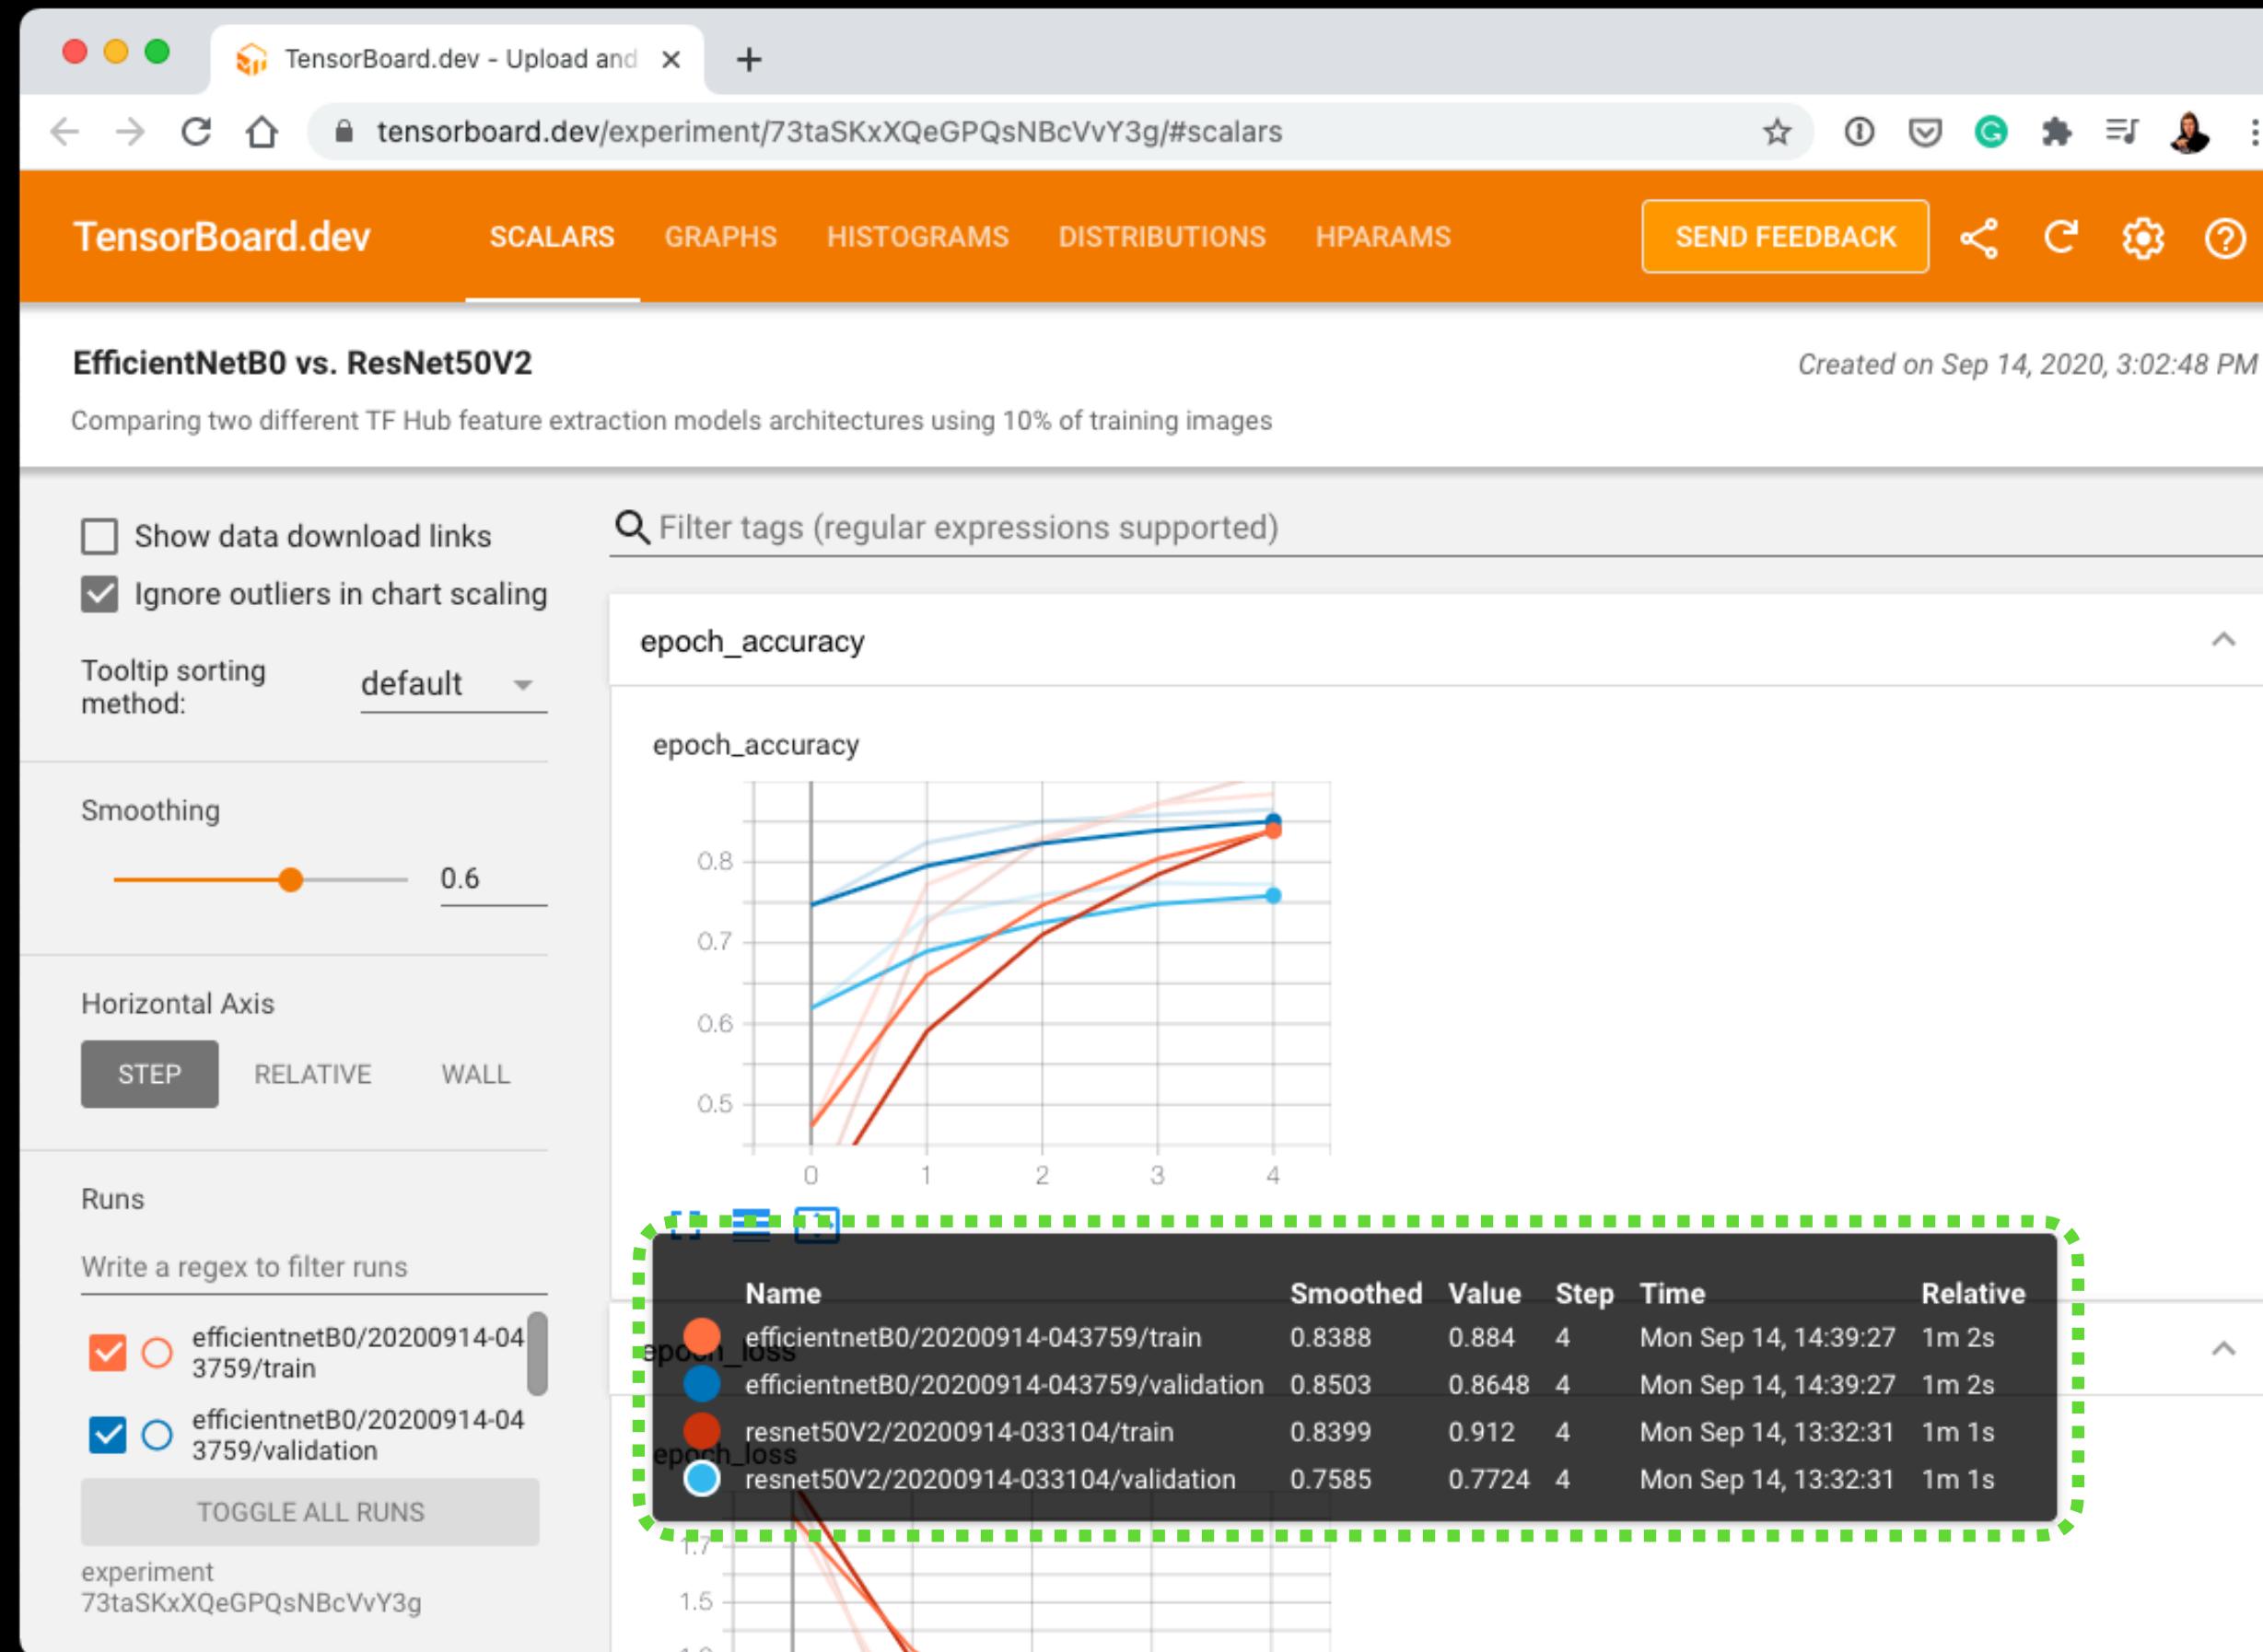
Fine-tuning

Kinds of Transfer Learning

Transfer Learning Type	Description	What happens	When to use
Original model (“As is”)	Take a pretrained model as it is and apply it to your task without any changes.	The original model remains unchanged.	Helpful if you have the exact same kind of data the original model was trained on.
Feature extraction	Take the underlying patterns (also called weights) a pretrained model has learned and adjust its outputs to be more suited to your problem.	Most of the layers in the original model remain frozen during training (only the top 1-3 layers get updated).	Helpful if you have a small amount of custom data (similar to what the original model was trained on) and want to utilise a pretrained model to get better results on your specific problem.
Fine-tuning	Take the weights of a pretrained model and adjust (fine-tune) them to your own problem.	Some (1-3+), many or all of the layers in the pretrained model are updated during training.	Helpful if you have a large amount of custom data and want to utilise a pretrained model and improve its underlying patterns to your specific problem.

What is TensorBoard?

- A way to **visually explore** your machine learning models performance and internals
- Host, track and share your machine learning experiments on TensorBoard.dev



(TensorBoard also integrates
with websites like Weights & Biases)

Comparing the results of two different model architectures (ResNet50V2 & EfficientNetB0) on the same dataset.

Source: <https://tensorboard.dev/experiment/73taSKxXQeGPQsNBcVvY3g/#scalars>